



LFD259

Kubernetes for Developers

Version 1.0



LFD259: Version 1.0

Copyright The Linux Foundation 2018. All rights reserved.

Copyright The Linux Foundation 2018. All rights reserved.

The training materials provided or developed by The Linux Foundation in connection with the training services are protected by copyright and other intellectual property rights.

Open source code incorporated herein may have other copyright holders and is used pursuant to the applicable open source license.

The training materials are provided for individual use by participants in the form in which they are provided. They may not be copied, modified, distributed to non-participants or used to provide training to others without the prior written consent of The Linux Foundation.

No part of this publication may be reproduced, photocopied, stored on a retrieval system, or transmitted without express prior written consent.

Published by:

the **Linux Foundation**

<http://www.linuxfoundation.org>

No representations or warranties are made with respect to the contents or use of this material, and any express or implied warranties of merchantability or fitness for any particular purpose or specifically disclaimed.

Although third-party application software packages may be referenced herein, this is for demonstration purposes only and shall not constitute an endorsement of any of these software applications.

Linux is a registered trademark of Linus Torvalds. Other trademarks within this course material are the property of their respective owners.

If there are any questions about proper and fair use of the material herein, please contact:

training@linuxfoundation.org



Exercise 1.1: Obtaining Some System Information

- Open up a command line terminal window; there are a number of ways to do this as we shall discuss later. If you do not see it under Applications->Accessories or Applications->System Tools, do Alt-F2 and type in **gnome-terminal**.
- The **uname** utility will provide some basics. If you type

```
$ uname --help
```

you will get:

```
$ uname --help
Usage: uname [OPTION]...
Print certain system information.  With no OPTION, same as -s.

-a, --all                print all information, in the following order,
                        except omit -p and -i if unknown:
-s, --kernel-name        print the kernel name
-n, --nodename           print the network node hostname
-r, --kernel-release     print the kernel release
-v, --kernel-version     print the kernel version
-m, --machine            print the machine hardware name
-p, --processor          print the processor type or "unknown"
-i, --hardware-platform  print the hardware platform or "unknown"
-o, --operating-system    print the operating system
    --help              display this help and exit
    --version            output version information and exit
```

Try some of the options.

- Other basic commands you might try are:

```
$ df
$ free
$ less /proc/meminfo
$ more /proc/cpuinfo
$ cat /proc/version
```



Lab 2.1 - Deploying a New Cluster

Overview

We will create a two-node Ubuntu 16.04 cluster. Using two nodes allows an understanding of issues and configurations found in a production environment. While 2 vCPU and 8G of memory allows for quick labs, you could also use much smaller VMs. Other Linux distributions should work in a very similar manner, but have not been tested.

Regardless of the platform used, VirtualBox, VMWare, AWS, GCE, or even bare metal, please remember that security software like SELinux and Firewalls can prevent the labs from working. While this is not something you should do in production, consider disabling the firewall and security software. GCE requires a new VPC to be created and a rule allowing all traffic to be included. The use of Wireshark can be a helpful place to start with troubleshooting, if you're unable to open all ports. Currently, **kubeadm** requires that swap to be turned off on every node. The **swapoff -a** command will do this until your next reboot, with various methods to disable swap persistently. Cloud providers typically provide instances with swap disabled.

To assist with setting up your cluster, please download the tarball of shell scripts and YAML files. The **k8sMaster.sh** and **k8sSecond.sh** scripts deploy a Kubernetes cluster using **kubeadm** and use **Project Calico** for networking.

```
student@ckad-1:~$ wget https://training.linuxfoundation.org/cm/LFD259/
--user=LFtraining --password=Penguin2014
student@ckad-1:~$ tar -xvf lfd259-example-files.tar
```

Deploy a Master Node using kubeadm

Review the script to install and begin the configuration of the master Kubernetes server.

```
student@ckad-1:~$ cat lfd259/k8sMaster.sh
#!/bin/bash -x
echo "This script is written to work with Ubuntu 16.04"
sleep 3
echo
echo "Disable swap until next reboot"
echo
sudo swapoff -a

echo "Update the local node"
sudo apt-get update && sudo apt-get upgrade -y
echo
echo "Install Docker"
sleep 3

sudo apt-get install -y docker.io
echo
echo "Install kubeadm and kubectl"
sleep 3

sudo sh -c "echo 'deb http://apt.kubernetes.io/ kubernetes-xenial main' >>
/etc/apt/sources.list.d/kubernetes.list"
```

<output_omitted>

Run the script as an argument to the bash shell. You will need the **kubeadm join** command shown near the end of the output when you add the minion node in a future step.

```
student@ckad-1:~$ bash lfd259/k8sMaster.sh
<output_omitted>
```

Your Kubernetes master has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

You should now deploy a pod network to the cluster.

Run `kubectl apply -f [podnetwork].yaml` with one of the options listed at:

<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

You can now join any number of machines by running the following on each node as root:

```
kubeadm join --token 118c3e.83b49999dc5dc034 10.128.0.3:6443
--discovery-token-ca-cert-hash
sha256:40aa946e3f53e38271bae24723866f56c86d77efb49aedeb8a70cc189bfe2e1d
--2018-03-30 19:32:41-- https://goo.gl/eWLkzb
```

<output_omitted>

NAME	STATUS	ROLES	AGE	VERSION
ckad-1	Ready	master	1m	v1.9.1

Deploy a Minion Node

Open a separate terminal into your second node. Having both terminal sessions allows you to monitor the status of the cluster while adding the second node.

```
student@ckad-2:~$ cat lfd259/k8sSecond.sh
#!/bin/bash -x
sudo apt-get update && sudo apt-get upgrade -y
```

<output_omitted>

Run the script on the second node:

```
student@ckad-2:~$ bash lfd259/k8sSecond.sh
```

<output_omitted>

When the script is done, the minion node is ready to join the cluster. The `kubeadm join` statement can be found near the end of the `kubeadm init` output. Your nodes will use a different IP address

and hashes than the example below. You will need to prepend **sudo** to run the script copied from the master node.

```
student@ckad-2:~$ sudo kubeadm join --token 118c3e.83b49999dc5dc034 \
  10.128.0.3:6443 --discovery-token-ca-cert-hash \
  sha256:40aa946e3f53e38271bae24723866f56c86d77efb49aedeb8a70cc189bfe2e1d
```

Configure the Master Node

Return to the master node. We will configure command line completion and verify that both nodes have been added to the cluster. The first command will configure completion in the current shell. The second command will ensure future shells have completion:

```
student@ckad-1:~$ source <(kubectl completion bash)
student@ckad-1:~$ echo "source <(kubectl completion bash)" >> ~/.bashrc
```

Verify that both nodes are part of the cluster. It may take a minute for the second node to reach a *Ready* state.

```
student@ckad-1:~$ kubectl get node
```

NAME	STATUS	ROLES	AGE	VERSION
ckad-1	Ready	master	5m	v1.9.1
ckad-2	Ready	<none>	2m	v1.9.1

Create a Simple Deployment

We will use the **kubectl** command for the majority of work with Kubernetes. Review the help output to become familiar with commands options and arguments:

```
student@ckad-1:~$ kubectl --help
kubectl controls the Kubernetes cluster manager.
```

Find more information at: <https://kubernetes.io/docs/reference/kubectl/overview/>.

Basic Commands (Beginner):

```
create      Create a resource from a file or from stdin.
expose      Take a replication controller, service, deployment or pod
and
```

expose it as a new Kubernetes Service

```
run          Run a particular image on the cluster
set          Set specific features on objects
run-container Run a particular image on the cluster. This command is
deprecated, use "run" instead
```

Basic Commands (Intermediate):

<output_omitted>

With more than 40 arguments, you can explore each by also using the `--help` option. Take a closer look at a few, starting with *taint*, for example:

```
student@ckad-1:~$ kubectl taint --help
Update the taints on one or more nodes.
```

```
* A taint consists of a key, value, and effect. As an argument here, it
is
expressed as key=value:effect.
```

```
* The key must begin with a letter or number, and may contain letters,
numbers, hyphens, dots, and underscores, up to 253 characters.
```

```
* Optionally, the key can begin with a DNS subdomain prefix and a single
'/',
like example.com/my-app
<output_omitted>
```

By default, the master node will not allow general containers to be deployed for security reasons. This is via a *taint*. Only containers which tolerate this taint will be scheduled on this node. As we only have two nodes in our cluster, we will remove the taint, allowing containers to be deployed on both nodes. The following command will remove the taint from all nodes, so you should see one success and one “not found” error. The minion node does not have the taint to begin with. Note the minus sign at the end of the command, which removes the preceding value.

```
student@ckad-1:~$ kubectl taint nodes --all node-role.kubernetes.io/master-
node "ckad-1" untainted
taint "node-role.kubernetes.io/master:" not found
```

Now, run a containerized webserver **nginx**. Use `kubectl run` to create a simple, single replica deployment running the **nginx** web server.

```
student@ckad-1:~$ kubectl run firstpod --image=nginx
deployment.apps "firstpod" created
```


Verify the new deployment exists and that the desired number of Pods matches the current number. Using a comma, you can request two resource types at once. The <Tab> key can be helpful. Type enough of the word to be unique and press the Tab key - it should complete the word. The deployment should show a number 1 for each value, such that the desired number of pods matches the up-to-date and running number. The pod should show zero restarts.

```
student@ckad-1:~$ kubectl get deployment,pod
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE
deployment.extension/firstpod	1	1	1	1

```
13s
```

NAME	READY	STATUS	RESTARTS	AGE
pod/firstpod-7d99ffc75-247kl	1/1	Running	0	13s

View the details of the deployment, then the pod. Work through the output slowly. Knowing what a healthy deployment and pod look like can be helpful when troubleshooting issues. Again, the <Tab> key can be helpful when using long auto-generated object names. You should be able to type `firstpod<Tab>` and the name will complete when viewing the pod.

```
student@ckad-1:~$ kubectl describe deployment firstpod
```

```
Name: firstpod
Namespace: default
CreationTimestamp: Fri, 30 Mar 2018 16:46:57 +0000
Labels: run=firstpod
Annotations: deployment.kubernetes.io/revision=1
Selector: run=firstpod
Replicas: 1 desired | 1 updated | 1 total | 1 available | 0
unavailable
StrategyType: RollingUpdate
MinReadySeconds: 0
<output_omitted>
```

```
student@ckad-1:~$ kubectl describe pod firstpod-7d99ffc75-247kl
```

```
Name: firstpod-7d99ffc75-247kl
Namespace: default
Node: ckad-2/10.128.0.2
Start Time: Fri, 30 Mar 2018 16:46:57 +0000
Labels: pod-template-hash=385599731
run=firstpod
Annotations: <none>
Status: Running
IP: 192.168.55.100
Controlled By: ReplicaSet/firstpod-7d99ffc75
```

```
Containers:
  firstpod:
    <output_omitted>
```

Note that the resources are in the `default` namespace. Get a list of available namespaces:

```
student@ckad-1:~$ kubectl get namespaces
NAME          STATUS    AGE
default       Active    20m
kube-public   Active    20m
kube-system   Active    20m
```

There are two other namespaces. Look at the pods in the `kube-system` namespace:

```
student@ckad-1:~$ kubectl get pod -n kube-system
NAME                                READY    STATUS    RESTARTS
AGE
calico-etcd-rvrpk                  1/1      Running    1
20m
calico-kube-controllers-d554689d5-lm687  1/1      Running    1
20m
calico-node-2ck9g                  2/2      Running    4
19m
calico-node-kkxvl                  2/2      Running    3
20m
etcd-ckad-1                        1/1      Running    1
20m
<output_omitted>
```

Now, look at the pods in a namespace that does not exist. Note that you do not receive an error:

```
student@ckad-1:~$ kubectl get pod -n fakenamepace
No resources found.
```

You can also view resources in all namespaces at once:

```
student@ckad-1:~$ kubectl get pod --all-namespaces
NAMESPACE    NAME                                READY    STATUS
RESTARTS    AGE
default      firstpod-7d99ffc75-247kl          1/1      Running
0            5m
```

```

kube-system    calico-etcd-z49kx                                1/1    Running
0              31m
kube-system    calico-kube-controllers-d554689d5-pfszw        1/1    Running
0              31m
<output_omitted>

```

View several resources at once. Note that most resources have a short name, such as **rs** for ReplicaSet, **po** for Pod, **svc** for Service, and **ep** for endpoint.

```

student@ckad-1:~$ kubectl get deploy,rs,po,svc,ep
NAME                                     DESIRED    CURRENT    UP-TO-DATE    AVAILABLE
AGE
deployment.extensions/firstpod          1          1          1             1
4m

NAME                                     DESIRED    CURRENT    READY    AGE
replicaset.extensions/firstpod-7d99ffc75 1          1          1        4m

NAME             READY    STATUS    RESTARTS    AGE
pod/firstpod-7d99ffc75 1/1      Running    0           4m

NAME             TYPE             CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
service/kubernetes ClusterIP        10.96.0.1     <none>         443/TCP    21m

NAME             ENDPOINTS            AGE
endpoints/kubernetes 10.128.0.3:6443    21m

```

Delete the *ReplicaSet* and view the resources again. Note that the time on the *ReplicaSet* and the Pod it controls is now less than a minute. The deployment controller restarted the *ReplicaSet*, which restarted the Pod when the desired configuration did not match the current status.

```

student@ckad-1:~$ kubectl delete rs firstpod-7d99ffc75
replicaset.extensions "firstpod-7d99ffc75" deleted

```

```

student@ckad-1:~$ kubectl get deployment,rs,po,svc,ep
NAME                                     DESIRED    CURRENT    UP-TO-DATE    AVAILABLE
AGE
deployment.extensions/firstpod          1          1          1             1
7m

```

NAME	DESIRED	CURRENT	READY
AGE			
replicaset.extensions/firstpod-7d99ffc75	1	1	1
12s			

NAME	READY	STATUS	RESTARTS	AGE
pod/firstpod-7d99ffc75-p9hbw	1/1	Running	0	12s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	24m

NAME	ENDPOINTS	AGE
endpoints/kubernetes	10.128.0.2:6443	24m

This time, delete the top-level controller. After about 30 seconds for everything to shut down, you should only see the cluster service and endpoint remain:

```
student@ckad-1:~$ kubectl delete deployment firstpod
deployment.extensions "firstpod" deleted
```

```
student@ckad-1:~$ kubectl get deployment,rs,po,svc,ep
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	24m

NAME	ENDPOINTS	AGE
kubernetes	10.128.0.3:6443	24m



Lab 3.1 - Deploying a New Application

Overview

In this lab, we will deploy a very simple Python application, test it using Docker, ingest it into Kubernetes, and configure probes to ensure it continues to run. This lab requires the completion of the previous lab, the installation and configuration of a Kubernetes cluster.

Working with Python

Install Python on your master node. It may already be installed, as is shown in the output below:

```
student@ckad-1:~$ sudo apt-get -y install python
Reading package lists... Done
Building dependency tree
Reading state information... Done
python is already the newest version (2.7.12-1~16.04).
python set to manually installed.
0 upgraded, 0 newly installed, 0 to remove and 5 not upgraded.
student@ckad-1:~$
```

Locate the Python binary on your system:

```
student@ckad-1:~$ which python
/usr/bin/python
```

Create a new directory and change into it. The `docker build` process pulls everything from the current directory into the image file by default. Make sure the chosen directory is empty:

```
student@ckad-1:~$ mkdir appl
```

```
student@ckad-1:~$ cd appl
```

```
student@ckad-1:~/appl$ ls -l
total 0
```

Create a simple Python script which prints the time and hostname every 5 seconds. There are six commented parts to this script, which should explain what each part is meant to do. The script is included with others in the course tar file, though you are encouraged to create the file by hand if you are not already familiar with the process:

```
student@ckad-1:~/appl$ vim simple.py
#!/usr/bin/python
## Import the necessary modules
import time
import socket

## Use an ongoing while loop to generate output
while True :

## Set the hostname and the current date
    host = socket.gethostname()
    date = time.strftime("%Y-%m-%d %H:%M:%S")

## Convert the date output to a string
    now = str(date)

## Open the file named date in append mode
## Append the output of hostname and time
    f = open("date.out", "a" )
    f.write(now + "\n")
    f.write(host + "\n")
    f.close()

## Sleep for five seconds then continue the loop
    time.sleep(5)
```

Make the file executable and test that it works. Use `<ctrl-c>` to interrupt the `while` loop after 20 or 30 seconds. The output will be sent to a newly created file in your current directory called `date.out`.

```
student@ckad-1:~/appl$ chmod +x simple.py
```

```

student@ckad-1:~/app1$ ./simple.py
^CTraceback (most recent call last):
  File "./simple.py", line 42, in <module>
    time.sleep(5)
KeyboardInterrupt

```

View the `date.out` file. It should contain the hostname and timestamp stamps:

```

student@ckad-1:~/app1$ cat date.out
2018-03-22 15:51:38
ckad-1
2018-03-22 15:51:43
ckad-1
2018-03-22 15:51:48
ckad-1
<output_omitted>

```

Create a Dockerfile. Note the name is important; it cannot have a suffix. We will use three statements: **FROM** to declare which version of Python to use, **ADD** to include our script, and **CMD** to indicate the action of the container. Should you be including more complex tasks, you may need to install extra libraries, shown commented out as **RUN pip install** in the following example:

```

student@ckad-1:~/app1$ vim Dockerfile
FROM python:2
ADD simple.py /
## RUN pip install pystrich
CMD [ "python", "./simple.py" ]

```

Build the container. The output below shows mid-build, as necessary software is downloaded. You will need to use **sudo** in order to run this command. After the three-step process completes, the last line of output should indicate success.

```

student@ckad-1:~/app1$ sudo docker build -t simpleapp .
Sending build context to Docker daemon 3.072 kB
Step 1/3 : FROM python:2
2: Pulling from library/python
4176fe04cefe: Pull complete
851356ecf618: Pull complete
6115379c7b49: Pull complete
aaf7d781d601: Extracting [=====]
54.03 MB/135 MB
40cf661a3cc4: Download complete
]

```



```
c582f0b73e63: Download complete
6c1ea8f72a0d: Download complete
7051a41ae6b7: Download complete
<output_omitted>
Successfully built c4e0679b9c36
```

Verify you can see the new image among others downloaded during the build process, installed to support the cluster, or you may have already worked with. The newly created `simpleapp` image should be listed first:

```
student@ckad-1:~/app1$ sudo docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
simpleapp	latest	c4e0679b9c36	2 minutes ago	681 MB
quay.io/calico/node	v2.6.8	e96a297310fd	13 days ago	282 MB
python	2	d8690ef56706	2 weeks ago	681 MB

```
<output_omitted>
```

Use Docker to run a container using the new image. While the script is running, you won't see any output and the shell will be occupied running the image in the background. After 30 seconds, use `<ctrl>-c` to interrupt. The local `date.out` file will not be updated with new times; instead, that output will be a file of the container image.

```
student@ckad-1:~$ sudo docker run simpleapp
^CTraceback (most recent call last):
  File "./simple.py", line 24, in <module>
    time.sleep(5)
KeyboardInterrupt
```

Locate the newly created `date.out` file. The following command should show two files of this name, the one created when we ran `simple.py` and another under `/var/lib/docker` when run via a Docker container:

```
student@ckad-1:~/app1$ sudo find / -name date.out
/home/student/app1/date.out
/var/lib/docker/aufs/diff/ee814320c900bd24fad0c5db4a258d3c2b78a19cde629d7de7d27270d6a0c1f5/date.out
```

View the contents of the `date.out` file created via Docker. Note the need for `sudo`, as Docker created the file this time, and the owner is root. The long name is shown on several lines in the example, but would be a single line when typed or copied.

```
student@ckad-1:~/app1$ sudo tail \
/var/lib/docker/aufs/diff/ee814320c900bd24fa\
d0c5db4a258d3c2b78a19cde629d7de7d27270d6a0c1f5/date.out
2018-03-22 16:13:46
53e1093e5d39
2018-03-22 16:13:51
53e1093e5d39
2018-03-22 16:13:56
53e1093e5d39
```

Configure A Local Docker Repository

While we could create an account and upload our application to `hub.docker.com`, thus sharing it with the world, we will instead create a local repository and make it available to the nodes of our cluster.

We'll need to complete a few steps with special permissions; for ease of use, we'll become root using `sudo`:

```
student@ckad-1:~/app1$ cd
student@ckad-1:~$ sudo -i
```

Install the `docker-compose` software and utilities to work with the `nginx` server, which will be deployed with the registry:

```
root@ckad-1:~# apt-get install -y docker-compose apache2-utils
<output_omitted>
```

Create a new directory for configuration information. We'll be placing the repository in the root filesystem. A better location may be chosen in a production environment.

```
root@ckad-1:~# mkdir -p /localdocker/data

root@ckad-1:~# cd /localdocker/
```

Create a `docker-compose` file. Inside is an entry for the `nginx` web server to handle outside traffic, and a `registry` entry listening to loopback port 5000 for running a local Docker registry.

```
root@ckad-1:/localdocker# vim docker-compose.yaml
nginx:
  image: "nginx:1.12"
  ports:
    - 443:443
  links:
    - registry:registry
  volumes:
    - /localdocker/nginx:/etc/nginx/conf.d
registry:
  image: registry:2
  ports:
    - 127.0.0.1:5000:5000
  environment:
    REGISTRY_STORAGE_FILESYSTEM_ROOTDIRECTORY: /data
  volumes:
    - /localdocker/data:/data
```

Use the `docker-compose up` command to create the containers declared in the previous step YAML file. This will capture the terminal and run until you use `<ctrl>-c` to interrupt. There should be five `registry_1` entries with `info` messages about memory and which port is being listened to. Once we're sure the `docker` file works, we'll convert to a Kubernetes tool.

```
root@ckad-1:/localdocker# docker-compose up
Pulling nginx (nginx:1.12)...
1.12: Pulling from library/nginx
2a72cbf407d6: Pull complete
f37cbdc183b2: Pull complete
78b5ad0b466c: Pull complete
Digest:
sha256:edad623fc7210111e8803b4359ba4854e101bccalfe7f46bd1d35781f4034f0c
Status: Downloaded newer image for nginx:1.12
Creating localdocker_registry_1
Creating localdocker_nginx_1
Attaching to localdocker_registry_1, localdocker_nginx_1
registry_1 | time="2018-03-22T18:32:37Z" level=warning msg="No HTTP secret
provided - generated ran
<output_omitted>
```

Test that you can access the repository. Open a second terminal to the master node. Use the `curl` command to test the repository. It should return `{}`, but does not have a carriage-return, so will be on the same line as the following prompt. You should also see the `GET` request in the first captured terminal, without error. Don't forget the trailing slash. You'll see a "**Moved Permanently**" message if the path doesn't match exactly.

```
student@ckad-1:~/localdocker$ curl http://127.0.0.1:5000/v2/
{}student@ckad-1:~/localdocker$
```

Now that we know that `docker-compose` format is working, ingest the file into Kubernetes using `kompose`. Use `<ctrl-c>` to stop the previous `docker-compose`.

```
^CGracefully stopping... (press Ctrl+C again to force)
Stopping localdocker_nginx_1 ... done
Stopping localdocker_registry_1 ... done
```

Download the `kompose` binary and make it executable:

```
root@ckad-1:/localdocker# curl -L
https://github.com/kubernetes/kompose/releases/download/v1.1.0/kompose-linux-amd64 -o kompose
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time
Current			Dload Upload	Total	Spent	Left
Speed						
100 609	0 609	0 0	1963 0	--:--:--	--:--:--	--:--:--
1970						
100 45.3M	100 45.3M	0 0	16.3M 0	0:00:02	0:00:02	--:--:--
25.9M						

```
root@ckad-1:/localdocker# chmod +x kompose
```

Move the binary to a directory in our `$PATH`. Then, return to your non-root user:

```
root@ckad-1:/localdocker# mv ./kompose /usr/local/bin/kompose
```

```
root@ckad-1:/localdocker# exit
```

Create two physical volumes in order to deploy a local registry for Kubernetes. 200mi for each should be enough for each of the volumes. More details on how persistent volumes and persistent volume claims are covered in an upcoming chapter.

```
student@ckad-1:~$ vim vol1.yaml
apiVersion: v1
kind: PersistentVolume
metadata:
  labels:
    type: local
  name: task-pv-volume
spec:
  accessModes:
  - ReadWriteOnce
  capacity:
    storage: 200Mi
  hostPath:
    path: /tmp/data
  persistentVolumeReclaimPolicy: Retain
```

```
student@ckad-1:~$ vim vol2.yaml
apiVersion: v1
kind: PersistentVolume
metadata:
  labels:
    type: local
  name: registryvm
spec:
  accessModes:
  - ReadWriteOnce
  capacity:
    storage: 200Mi
  hostPath:
    path: /tmp/nginx
  persistentVolumeReclaimPolicy: Retain
```

Create both volumes:

```
student@ckad-1:~$ kubectl create -f vol1.yaml
persistentvolume "task-pv-volume" created
```

```
student@ckad-1:~$ kubectl create -f vol2.yaml
persistentvolume "registryvm" created
```

Verify that both volumes have been created. They should show an **Available** status:

```
student@ckad-1:~$ kubectl get pv
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS
registryvm	200Mi	RWO	Retain	Available
task-pv-volume	200Mi	RWO	Retain	Available

Go to the configuration file for a `localdocker` registry:

```
student@ckad-1:~$ cd /localdocker/
```

```
student@ckad-1:~/localdocker$ ls
data  docker-compose.yaml  nginx
```

Convert the Docker file into a single YAML file for use with Kubernetes. Not all objects convert exactly from Docker to kompose; you will get errors about the mount syntax for the new volumes. They can be safely ignored.

```
student@ckad-1:~/localdocker$ sudo kompose convert -f docker-compose.yaml \
-o localregistry.yaml
WARN Volume mount on the host "/localdocker/nginx/" isn't supported -
ignoring path on the host
WARN Volume mount on the host "/localdocker/data" isn't supported -
ignoring path on the host
```

Review the file. You'll find that multiple objects will be created as well:

```
student@ckad-1:~/localdocker$ less localregistry.yaml
apiVersion: v1
items:
- apiVersion: v1
  kind: Service
  metadata:
    annotations:
      kompose.cmd: kompose convert -f docker-compose.yaml -o
localregistry.yaml
      kompose.version: 1.1.0 (36652f6)
      creationTimestamp: null
    labels:
  <output_omitted>
```

View the cluster resources prior to deploying the registry. Only the cluster service and two available persistent volumes should exist in the `default` namespace:

```
student@ckad-1:~/localdocker$ kubectl get pods,svc,pvc,pv,deploy
NAME                                TYPE            CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
kubernetes                          ClusterIP       10.96.0.1     <none>         443/TCP    4h
```



```
NAME                                CAPACITY    ACCESS MODES    RECLAIM POLICY    STATUS
CLAIM    STORAGECLASS    REASON    AGE
registryvm    200Mi    RWO                Retain            Available
15s
task-pv-volume    200Mi    RWO                Retain            Available
17s
```

Use `kubectl` to create a local Docker registry:

```
student@ckad-1:~/localdocker$ kubectl create -f localregistry.yaml
service "nginx" created
service "registry" created
deployment.extensions "nginx" created
persistentvolumeclaim "nginx-claim0" created
deployment.extensions "registry" created
persistentvolumeclaim "registry-claim0" created
```

View the newly deployed resources. The persistent volumes should now show as **Bound**. Find the service IP for the registry. It should be sharing port 5000. In the example below, the IP address is 10.110.186.162, but yours may be different:

```
student@ckad-1:~/localdocker$ kubectl get pods,svc,pvc,pv,deploy
NAME                                READY    STATUS    RESTARTS    AGE
pod/nginx-6b58d9cdfd-95zxq          1/1      Running    0            1m
pod/registry-795c6c8b8f-b8z4k       1/1      Running    0            1m
```



```
NAME                                TYPE            CLUSTER-IP    EXTERNAL-IP    PORT(S)
AGE
service/kubernetes                  ClusterIP       10.96.0.1     <none>         443/TCP
1h
service/nginx                        ClusterIP       10.106.82.218 <none>         443/TCP
1m
service/registry                    ClusterIP       10.110.186.162 <none>         5000/TCP
1m
```

NAME	STATUS	VOLUME	CAPACITY
ACCESS MODES STORAGECLASS AGE			
persistentvolumeclaim/nginx-claim0	Bound	registryvm	200Mi
RWO			1m
persistentvolumeclaim/registry-claim0	Bound	task-pv-volume	200Mi
RWO			1m

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY
STATUS CLAIM		REASON	AGE
persistentvolume/registryvm	200Mi	RWO	Retain
Bound default/nginx-claim0			5m
persistentvolume/task-pv-volume	200Mi	RWO	Retain
Bound default/registry-claim0			6m

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE
AGE				
deployment.extensions/nginx	1	1	1	1
1m				
deployment.extensions/registry	1	1	1	1
1m				

Verify you get the same { } response using the Kubernetes deployed registry, as we did when using `docker-compose`.

Note you must use the trailing slash after `v2`. Please also note that, if the connection hangs, it may be due to a firewall issue. If running your nodes using GCE, ensure your instances are using VPC setup and all ports are allowed. If using AWS, also make sure all ports are being allowed.

```
student@ckad-1:~/localdocker$ curl http://10.110.186.162:5000/v2/
{}student@ckad-1:~/localdocker$
```

Edit the Docker configuration file to allow insecure access to the registry. In a production environment, steps should be taken to create and use TLS authentication instead. Use the IP and port of the registry:

```
student@ckad-1:~$ sudo vim /etc/docker/daemon.json
{ "insecure-registries": ["10.110.186.162:5000"] }
```

Restart Docker on the local system. It can take up to a minute for the restart to take place:

```
student@ckad-1:~$ sudo systemctl restart docker.service
```


Download and tag a typical image from hub.docker.com. Tag the image using the IP and port of the registry. We will also use the `latest` tag.

```
student@ckad-1:~$ sudo docker pull ubuntu
Using default tag: latest
latest: Pulling from library/ubuntu
<output_omitted>
Digest:
sha256:9ee3b83bcaa383e5e3b657f042f4034c92cdd50c03f73166c145c9ceaea9ba7c
Status: Downloaded newer image for ubuntu:latest

student@ckad-1:~$ sudo docker tag ubuntu:latest 10.110.186.162:5000/tagtest
```

Push the newly tagged image to your local registry. If you receive an error about an HTTP request to an HTTPS client, check that you edited the `/etc/docker/daemon.json` file correctly and restarted the service:

```
student@ckad-1:~$ sudo docker push 10.110.186.162:5000/tagtest
The push refers to a repository [10.110.186.162:5000/tagtest]
db584c622b50: Pushed
52a7ea2bb533: Pushed
52f389ea437e: Pushed
88888b9b1b5b: Pushed
a94e0d5a7c40: Pushed
latest: digest:
sha256:0847cc7fed1bfafac713b0aa4ddfb8b9199a99092ae1fc4e718cb28e8528f65f
size: 1357
```

We will test to make sure we can also pull images from our local repository. Begin by removing the local cached images:

```
student@ckad-1:~$ sudo docker image remove ubuntu:latest
Untagged: ubuntu:latest
Untagged:
ubuntu@sha256:e348fbbbea0e0a0e73ab0370de151e7800684445c509d46195aef73e090a49bd6

student@ckad-1:~$ sudo docker image remove 10.110.186.162:5000/tagtest
Untagged: 10.110.186.162:5000/tagtest:latest
Untagged:
10.110.186.162:5000/tagtest@sha256:0847cc7fed1bfafac713b0aa4ddfb8b9199a99092ae1fc4e718cb28e8528f65f
```

```
Deleted:
sha256:f975c50357489439eb9145dbfa16bb7cd06c02c31aa4df45c77de4d2baa4e232
Deleted:
sha256:0bd983fc698ee9453dd7d21f8572ea1016ec9255346ceabb0f9e173b4348644f
Deleted:
sha256:08fe90e1a1644431accc00cc80f519f4628dbf06a653c76800b116d3333d2b6d
Deleted:
sha256:5dc5eef2b94edd185b4d39586e7beb385a54b6bac05d165c9d47494492448235
Deleted:
sha256:14a40a140881d18382e13b37588b3aa70097bb4f3fb44085bc95663bdc68fe20
Deleted:
sha256:a94e0d5a7c404d0e6fa15d8cd4010e69663bd8813b5117fbad71365a73656df9
```

Pull the image from the local registry. It should report the download of a newer image:

```
student@ckad-1:~$ sudo docker pull 10.110.186.162:5000/tagtest
Using default tag: latest
latest: Pulling from tagtest
bf8f2f4f7b8b: Pull complete
4288a6810024: Pull complete
bc5512367466: Pull complete
aa9fdb4f8e2a: Pull complete
665607941289: Pull complete
Digest:
sha256:0847cc7fed1bfafac713b0aa4ddfb8b9199a99092ae1fc4e718cb28e8528f65f
Status: Downloaded newer image for 10.110.186.162:5000/tagtest:latest
```

Use `docker tag` to assign the `simpleapp` image, and then `push` it to the local registry. The image and dependent images should be pushed to the local repository:

```
student@ckad-1:~$ sudo docker tag simpleapp 10.110.186.162:5000/simpleapp

student@ckad-1:~$ sudo docker push 10.110.186.162:5000/simpleapp
The push refers to a repository [10.110.186.162:5000/simpleapp]
321938b97e7e: Pushed
ca82a2274c57: Pushed
de2fbb43bd2a: Pushed
4e32c2de91a6: Pushed
6e1b48dc2ccc: Pushed
ff57bdb79ac8: Pushed
6e5e20cbf4a7: Pushed
86985c679800: Pushed
8fad67424c4e: Pushed
```

```
latest: digest:
sha256:67ea3e11570042e70cdcbad684a1e2986f59aaf53703e51725accdf5c70d475a
size: 2218
```

Configure the second minion node to use the local registry running on the master server. Connect to the minion node. Edit the Docker file with the same values from the master node, and restart the service:

```
student@ckad-2:~$ sudo vim /etc/docker/daemon.json
{ "insecure-registries":["10.110.186.162:5000"] }
```

```
student@ckad-2:~$ sudo systemctl restart docker.service
```

Pull the recently pushed image from the registry running on the master node:

```
student@ckad-2:~$ sudo docker pull 10.110.186.162:5000/simpleapp
Using default tag: latest
latest: Pulling from simpleapp
f65523718fc5: Pull complete
1d2dd88bf649: Pull complete
c09558828658: Pull complete
0e1d7c9e6c06: Pull complete
c6b6fe164861: Pull complete
45097146116f: Pull complete
f21f8abae4c4: Pull complete
1c39556edcd0: Pull complete
85c79f0780fa: Pull complete
Digest:
sha256:67ea3e11570042e70cdcbad684a1e2986f59aaf53703e51725accdf5c70d475a
Status: Downloaded newer image for 10.110.186.162:5000/simpleapp:latest
```

Return to the master node and deploy the **simpleapp** in **kubernetes** with several replicas. We will name the deployment **try1**. With multiple replicas, the scheduler should run some containers on each node:

```
student@ckad-1:~$ kubectl run try1 \
  --image=10.110.186.162:5000/simpleapp:latest \
  --replicas=6
deployment.apps "try1" created
```

View the running pods. You should see six replicas of **simpleapp**, as well as two running the locally hosted image repository:

```
student@ckad-1:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-j6jm6	1/1	Running	1	13m
registry-795c6c8b8f-5jnnp	1/1	Running	1	13m
try1-857bdcd888-6klrr	1/1	Running	0	25s
try1-857bdcd888-9pwnp	1/1	Running	0	25s
try1-857bdcd888-9xkth	1/1	Running	0	25s
try1-857bdcd888-tw58z	1/1	Running	0	25s
try1-857bdcd888-xj9lk	1/1	Running	0	25s
try1-857bdcd888-znpm8	1/1	Running	0	25s

On the second node, use `docker ps` to verify containers of `simpleapp` are running. The scheduler will try to deploy an equal number to both nodes:

```
student@ckad-2:~$ sudo docker ps | grep simple
```

```
3ae4668d71d8
10.110.186.162:5000/simpleapp@sha256:67ea3e11570042e70cdcbad684a1e2986f59aa
f53703e51725accdf5c70d475a          "python ./simple.py"
48 seconds ago          Up 48 seconds
k8s_try1_try1-857bdcd888-9xkth_default_2e94b97e-322a-11e8-af56-42010a800004
_0

ef6448764625
10.110.186.162:5000/simpleapp@sha256:67ea3e11570042e70cdcbad684a1e2986f59aa
f53703e51725accdf5c70d475a          "python ./simple.py"
48 seconds ago          Up 48 seconds
k8s_try1_try1-857bdcd888-znpm8_default_2e99f356-322a-11e8-af56-42010a800004
_0
```

Return to the master node. Save the `try1` deployment as YAML:

```
student@ckad-1:~/app1$ cd ~/app1/
student@ckad-1:~/app1$ kubectl get deployment try1 -o yaml > simpleapp.yaml
```

Edit the YAML file to remove `creationTimestamp`, `selfLink`, `uid`, `resourceVersion`, and all the `status` information. In newer versions of Kubernetes it seems to no longer be necessary to remove these values in order to deploy again. Be aware that older versions would error if these values were found in the YAML file. For backwards compatibility, we will continue to remove these entries:

```
student@ckad-1:~/app1$ vim simpleapp.yaml
<output_omitted>
```

Delete and recreate the `try1` deployment using the YAML file. Verify the deployment is running with the expected number of replicas:

```
student@ckad-1:~$ kubectl delete deployment try1
deployment.extensions "try1" deleted
```

```
student@ckad-1:~/app1$ kubectl create -f simpleapp.yaml
deployment.extensions "try1" created
```

```
student@ckad-1:~/app1$ kubectl get deployment
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
nginx	1	1	1	1	17m
registry	1	1	1	1	17m
try1	6	6	6	6	7s

Configure Probes

When large datasets need to be loaded or a complex application launched prior to client access, a **readinessProbe** can be used. The Pod will not become available to the cluster until a test is met. **readinessProbes** and **livenessProbes** use the same syntax and are identical, other than the name. Where the **readinessProbe** is checked prior to being ready, then not again, the **livenessProbe** continues to be checked. There are three types of liveness probes:

- A command returns a zero exit value, meaning success
- An HTTP request returns a response code in the 200 to 500 range
- The third probe uses a TCP socket.

In this example, we'll use a command, `cat`, which will return a zero exit code when the file `/tmp/healthy` has been created and can be accessed.

Edit the YAML deployment file and add the stanza for a readiness probe. Remember that, when working with YAML, whitespace matters. Indentation is used to parse where information should be associated within the stanza and the entire file. If you get an error about validating data, check the indentation. It can also be helpful to paste the file to this website to see how indentation affects the JSON value, which is actually what Kubernetes ingests: <https://www.json2yaml.com/>:

```
student@ckad-1:~/app1$ vim simpleapp.yaml
...
spec:
  containers:
```

```

- image: 10.111.235.60:5000/simpleapp:latest
  imagePullPolicy: Always
  name: try1
  readinessProbe:
    exec:
      command:
      - cat
      - /tmp/healthy
    periodSeconds: 5
  resources: {}

```

....

Delete and recreate the `try1` deployment:

```

student@ckad-1:~/app1$ kubectl delete deployment try1
deployment.extensions "try1" deleted

```

```

student@ckad-1:~/app1$ kubectl create -f simpleapp.yaml
deployment.extensions "try1" created

```

The new `try1` deployment should reference six pods, but show zero available. They are all missing the `/tmp/healthy` file:

```

student@ckad-1:~/app1$ kubectl get deployment

```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
nginx	1	1	1	1	39m
registry	1	1	1	1	39m
try1	6	6	6	0	5s

Take a closer look at the pods. Choose one of the `try1` pods as a test to create the health check file:

```

student@ckad-1:~/app1$ kubectl get pods

```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-g7lnk	1/1	Running	1	40m
registry-795c6c8b8f-7vwdn	1/1	Running	1	40m
try1-9869bdb88-2wfnr	0/1	Running	0	26s
try1-9869bdb88-6bknl	0/1	Running	0	26s
try1-9869bdb88-786v8	0/1	Running	0	26s
try1-9869bdb88-gmvs4	0/1	Running	0	26s
try1-9869bdb88-lfvlx	0/1	Running	0	26s
try1-9869bdb88-rtchc	0/1	Running	0	26s

Run the bash shell interactively and touch the `/tmp/healthy` file:

```
student@ckad-1:~/app1$ kubectl exec -it try1-9869bdb88-rtchc -- /bin/bash
root@try1-9869bdb88-rtchc:/# touch /tmp/healthy
root@try1-9869bdb88-rtchc:/# exit
exit
```

Wait at least five seconds, then check the pods again. Once the probe runs again, the container should show available quickly. The pod with the existing `/tmp/healthy` file should be running and show 1/1 in a **READY** state. The rest will continue to show 0/1.

```
student@ckad-1:~/app1$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-g7lnk	1/1	Running	1	44m
registry-795c6c8b8f-7vwdn	1/1	Running	1	44m
try1-9869bdb88-2wfnr	0/1	Running	0	4m
try1-9869bdb88-6bkn1	0/1	Running	0	4m
try1-9869bdb88-786v8	0/1	Running	0	4m
try1-9869bdb88-gmvs4	0/1	Running	0	4m
try1-9869bdb88-lfv1x	0/1	Running	0	4m
try1-9869bdb88-rtchc	1/1	Running	0	4m

Touch the file in the remaining pods. Consider a for loop, as an easy method to update each pod:

```
student@ckad-1:~$ for name in try1-9869bdb88-2wfnr try1-9869bdb88-6bkn1
try1-9869bdb88-786v8 try1-9869bdb88-gmvs4 try1-9869bdb88-lfv1x
> do
> kubectl exec $name touch /tmp/healthy
> done
```

It may take a short while for the probes to check, for the file and the health checks to succeed:

```
student@ckad-1:~/app1$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-g7lnk	1/1	Running	1	1h
registry-795c6c8b8f-7vwdn	1/1	Running	1	1h
try1-9869bdb88-2wfnr	1/1	Running	0	22m
try1-9869bdb88-6bkn1	1/1	Running	0	22m
try1-9869bdb88-786v8	1/1	Running	0	22m
try1-9869bdb88-gmvs4	1/1	Running	0	22m

```
try1-9869bdb88-lfv1x      1/1      Running    0          22m
try1-9869bdb88-rtchc      1/1      Running    0          22m
```

Now that we know when a Pod is healthy, we may want to keep track that it stays healthy, using a `livenessProbe`. You could use one probe to determine when a Pod becomes available and a second probe, to a different location, to ensure ongoing health.

Edit the Deployment again. Add in a `livenessProbe` section as seen below. This time we will add a new container to the pod running a simple application which will respond to port 8080. Note that the dash (-) in front of the `name: goproxy` is indented the same amount as the - in front of the `image: simpleapp` earlier in the file. In this example that would be seven spaces:

```
student@ckad-1:~/app1$ vim simpleapp.yaml
...
    terminationMessagePath: /dev/termination-log
    terminationMessagePolicy: File
-   name: goproxy
    image: k8s.gcr.io/goproxy:0.1
    ports:
    -   containerPort: 8080
    readinessProbe:
        tcpSocket:
            port: 8080
        initialDelaySeconds: 5
        periodSeconds: 10
    livenessProbe:
        tcpSocket:
            port: 8080
        initialDelaySeconds: 15
        periodSeconds: 20
    dnsPolicy: ClusterFirst
    restartPolicy: Always
...

```

Delete and recreate the Deployment:

```
student@ckad-1:~$ kubectl delete deployment try1
deployment.extensions "try1" deleted

student@ckad-1:~$ kubectl create -f simpleapp.yaml
deployment.extensions "try1" created
```


View the newly created Pods. You'll note that there are two containers per pod, and only one is running. The new `simpleapp` containers will not have the `/tmp/healthy` file, so they will not become available until we touch the `/tmp/healthy` file again. We could include a command which creates the file into the container arguments. The output below shows it can take a bit for the old pods to terminate.

```
student@ckad-1:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-g7lnk	1/1	Running	1	13h
registry-795c6c8b8f-7vwdn	1/1	Running	1	13h
try1-76cc5ffcc6-4rjvh	1/2	Running	0	3s
try1-76cc5ffcc6-bk5f5	1/2	Running	0	3s
try1-76cc5ffcc6-d8n5q	0/2	ContainerCreating	0	3s
try1-76cc5ffcc6-mm6tw	1/2	Running	0	3s
try1-76cc5ffcc6-r9q5n	1/2	Running	0	3s
try1-76cc5ffcc6-tx4dz	1/2	Running	0	3s
try1-9869bdb88-2wfnr	1/1	Terminating	0	12h
try1-9869bdb88-6bkn1	1/1	Terminating	0	12h
try1-9869bdb88-786v8	1/1	Terminating	0	12h
try1-9869bdb88-gmvs4	1/1	Terminating	0	12h
try1-9869bdb88-lfv1x	1/1	Terminating	0	12h
try1-9869bdb88-rtchc	1/1	Terminating	0	12h

Create the health check file for the `readinessProbe`. You can use a for loop again for each action, with updated Pod names. As there are now two containers in the Pod, you should include the container name for where the command will execute. If no name is given, it will default to the first container. Depending on how you edited the YAML file, `try1` should be the first pod and `goproxy` the second. To ensure the correct container is updated, add `-c try1` to the `kubectl` command. Your Pod names will be different. Use the names of the newly started containers from the `kubectl get pods` command output.

```
student@ckad-1:~$ for name in try1-76cc5ffcc6-4rjvh try1-76cc5ffcc6-bk5f5
try1-76cc5ffcc6-d8n5q try1-76cc5ffcc6-mm6tw try1-76cc5ffcc6-r9q5n
try1-76cc5ffcc6-tx4dz
do
kubectl exec $name -c try1 touch /tmp/healthy
done
```

In the next minute or so, the second container in each Pod, which was not running, will change status to `Running`. Each should show 2/2 containers running:

```
student@ckad-1:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
------	-------	--------	----------	-----

nginx-6b58d9cdfd-g7lnk	1/1	Running	1	13h
registry-795c6c8b8f-7vwdn	1/1	Running	1	13h
try1-76cc5ffcc6-4rjvh	2/2	Running	0	3s
try1-76cc5ffcc6-bk5f5	2/2	Running	0	3s
try1-76cc5ffcc6-d8n5q	2/2	Running	0	3s
try1-76cc5ffcc6-mm6tw	2/2	Running	0	3s
try1-76cc5ffcc6-r9q5n	2/2	Running	0	3s
try1-76cc5ffcc6-tx4dz	2/2	Running	0	3s

View the events for a particular pod. Even though both containers are currently running and the pod is in good shape, note the events show the last issue:

```
student@ckad-1:~/app1$ kubectl describe pod try1-76cc5ffcc6-tx4dz | tail
Normal      SuccessfulMountVolume  9m              kubelet, ckad-1-lab-x6dj
MountVolume.SetUp succeeded for volume "default-token-jf69w"
Normal      Pulling                9m              kubelet,
ckad-1-lab-x6dj pulling image "10.108.143.90:5000/simpleapp"
Normal      Pulled                 9m              kubelet,
ckad-1-lab-x6dj Successfully pulled image "10.108.143.90:5000/simpleapp"
Normal      Created               9m              kubelet,
ckad-1-lab-x6dj Created container
Normal      Started               9m              kubelet,
ckad-1-lab-x6dj Started container
Normal      Pulling                9m              kubelet,
ckad-1-lab-x6dj pulling image "k8s.gcr.io/goproxy:0.1"
Normal      Pulled                 9m              kubelet,
ckad-1-lab-x6dj Successfully pulled image "k8s.gcr.io/goproxy:0.1"
Normal      Created               9m              kubelet,
ckad-1-lab-x6dj Created container
Normal      Started               9m              kubelet,
ckad-1-lab-x6dj Started container
Warning     Unhealthy              4m (x60 over 9m) kubelet,
ckad-1-lab-x6dj Readiness probe failed: cat: /tmp/healthy: No such file or
directory
```

If you look for the status of each container in the pod, they should show that both are running and ready.

```
student@ckad-1:~/app1$ kubectl describe pod try1-76cc5ffcc6-tx4dz | \
grep -E 'State|Ready'
State:      Running
Ready:      True
State:      Running
```

Ready:	True
Ready	True



Lab 4.1 - Planning the Deployment

Overview

In this exercise, we will investigate common network plugins. Each *kubelet* agent uses one plugin at a time. Due to complexity, the entire cluster uses one plugin which is configured prior to application deployment. Some plugins don't honor security configurations like network policies. Should you design a deployment which uses a network policy, there wouldn't be an error, you just would not be protected.

While still new, the community is moving towards the Container Network Interface (CNI) specification (<https://github.com/containernetworking/cni>). This provides the most flexibility and features. A common alternative is *kubenet*, a basic plugin which relies on the cloud provider to handle routing and cross-node networking. In a previous lab exercise, we configured Project Calico. Classic and external modes are also possible.

Evaluate Network Plugins

Verify your nodes are using a CNI plugin. Look for options passed to the *kubelet*. You may see other lines, including the **grep** command itself and a shell script running in a container which configures Calico.

```
student@ckad-2-nzjr:~$ ps -ef | grep cni
student  13473 13442  0 22:55 pts/1      00:00:00 grep --color=auto cni
root      14118      1  2 Mar30 ?          02:57:27 /usr/bin/kubelet
--bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf
--kubeconfig=/etc/kubernetes/kubelet.conf
--pod-manifest-path=/etc/kubernetes/manifests --allow-privileged=true
--network-plugin=cni --cni-conf-dir=/etc/cni/net.d
--cni-bin-dir=/opt/cni/bin --cluster-dns=10.96.0.10
```

```
--cluster-domain=cluster.local --authorization-mode=Webhook
--client-ca-file=/etc/kubernetes/pki/ca.crt --cadvisor-port=0
--rotate-certificates=true --cert-dir=/var/lib/kubelet/pki
root      30591 30570   0 Mar30 ?           00:00:10 /bin/sh /install-cni.sh
```

View the details of the `install-cni.sh` script. The script runs in a container, the path to which will be different than the example below. Read through the script to see what it does on our behalf:

```
student@ckad-2-nzjr:~$ sudo find / -name install-cni.sh
/var/lib/docker/aufs/mnt/e95a30499a76e79027502bbb8ee4eeb8464a657e276a493249
f573f5d86e19b3/install-cni.sh
/var/lib/docker/aufs/diff/a7cd14de39089493b793f135ec965f0f3f79eeec8f9d78e67
9d7be1a3bdf3345/install-cni.sh
```

```
student@ckad-2-nzjr:~$ sudo less
/var/lib/docker/aufs/diff/a7cd14de39089493b793f135ec965f0f3f79eeec8f9d78e67
9d7be1a3bdf3345/install-cni.sh
```

There are many CNI providers possible. The following list represents some of the more common choices, but it is not exhaustive. With many new plugins being developed, there may be another which better serves your needs. Use these websites to answer the questions which follow. While we strive to keep the answers accurate, please be aware that this area receives a lot of attention and development, and changes often.

- Project Calico <https://docs.projectcalico.org/v3.0/introduction/>
- Calico with Canal <https://docs.projectcalico.org/v3.0/getting-started/kubernetes/installation/hosted/canal/>
- Weave Works <https://www.weave.works/docs/net/latest/kubernetes/kube-addon>
- Flannel <https://github.com/coreos/flannel>
- Romana http://romana.io/how/romana_basics/
- Kube Router: <https://www.kube-router.io>
- Kopeio <https://github.com/kopeio/networking>

Which of the plugins allow vxlans?

Which are layer 2 plugins?

Which are layer 3?

Which allow network policies?

Which can encrypt all TCP and UDP traffic?

Answers:

Which of the plugins allow vxlans?

Canal, Flannel, Kopeio-networking and Weave Net.

Which are layer 2 plugins?

Canal, Flannel, Kopeio-networking and Weave Net.

Which are layer 3?

Project Calico, Romana, and Kube Router

Which allow network policies?

Project Calico, Canal, Kube Router, Romana and Weave Net

Which can encrypt all TCP and UDP traffic?

Project Calico, Kopeio, and weave Net.

Multi-Container Pod Considerations

Using the information learned from this chapter, consider the following questions:

Which deployment method would allow the most flexibility, multiple applications per pod or one per Pod?

Which deployment method allows for the most granular scalability?

Which have the best performance?

How many IP addresses are assigned per pod?

What are some ways containers can communicate within the same pod?

What are some reasons you should have multiple containers per pod?

Answers:

Which deployment method would allow the most flexibility, multiple applications per pod or one per Pod?

One per pod.

Which deployment method allows for the most granular scalability?

One per pod

Which have the best performance?

Multiple per pod.

How many IP addresses are assigned per pod?

One

What are some ways containers can communicate within the same pod?

IPC, loopback or shared filesystem access.

What are some reasons you should have multiple containers per pod?

Lean containers may not have functionality like logging. Able to maintain lean execution, but add functionality as necessary.



Lab 5.1 - Configuring the Deployment

Overview

In this lab, we will add resources to our deployment, with further configuration you may need for production. We'll also work with updating deployed applications and automation of batch jobs and regular tasks.

Save a copy of your `~/app1/simpleapp.yaml` file, in case you would like to repeat portions of the labs, or you find your file difficult to use due to typos and whitespace issues.

```
student@ckad-1:~$ cp ~/app1/simpleapp.yaml ~/beforeLab5.yaml
```

Secrets and ConfigMap

There are three different ways a ConfigMap can ingest data:

- From a literal value
- From a file
- From a directory of files.

Create a ConfigMap containing primary colors. We will create a series of files to ingest into the ConfigMap. First, create a directory **primary** and populate it with four files. Then, we create a file in our home directory with our favorite color:

```
student@ckad-1:~/app1$ cd
student@ckad-1:~$ mkdir primary
student@ckad-1:~$ echo c > primary/cyan
student@ckad-1:~$ echo m > primary/magenta
```



```
student@ckad-1:~$ echo y > primary/yellow
student@ckad-1:~$ echo k > primary/black
student@ckad-1:~$ echo "known as key" >> primary/black
student@ckad-1:~$ echo blue > favorite
```

Generate a `configmap` using each of the three methods:

```
student@ckad-1:~$ kubectl create configmap colors \
  --from-literal=text=black \
  --from-file=./favorite \
  --from-file=./primary/
configmap "colors" created
```

View the newly created `configmap`. Note the way the ingested data is presented:

```
student@ckad-1:~$ kubectl get configmap colors
NAME      DATA      AGE
colors    6          11s
```

```
student@ckad-1:~$ kubectl get configmap colors -o yaml
apiVersion: v1
data:
  black: |
    k
    known as key
  cyan: |
    c
  favorite: |
    blue
  magenta: |
    m
  text: black
  yellow: |
    y
kind: ConfigMap
metadata:
  creationTimestamp: 2018-04-05T19:49:59Z
  name: colors
  namespace: default
  resourceVersion: "13491"
  selfLink: /api/v1/namespaces/default/configmaps/colors
  uid: 86457ce3-390a-11e8-ba73-42010a800003
```

Update the YAML file of the application to make use of the `configmap` as an environmental parameter. Add the six lines from the `env:` line to `key: favorite`.

```
student@ckad-1:~$ vim appl/simpleapp.yaml
...
spec:
  containers:
  - image: 10.105.119.236:5000/simpleapp:latest
    env:
    - name: ilike
      valueFrom:
        configMapKeyRef:
          name: colors
          key: favorite
      imagePullPolicy: Always
...
```

Delete and re-create the deployment with the new parameters:

```
student@ckad-1-lab-7xtx:~$ kubectl delete deployment try1
deployment.extensions "try1" deleted

student@ckad-1-lab-7xtx:~$ kubectl create -f appl/simpleapp.yaml
deployment.extensions "try1" created
```

Even though the `try1` container is not in a `ready` state, it is running and useful. Use `kubectl exec` to view a variable's value. View the pod state, then verify you can see the `ilike` value within.

```
student@ckad-1:~$ kubectl get po
<output_omitted>
student@ckad-1:~$ kubectl exec -c try1 -it try1-5db9bc6f85-whxbf -- \
    /bin/bash -c 'echo $ilike'
blue
```

Edit the YAML file again, this time adding the third method of using a ConfigMap. Edit the file to add three lines. `envFrom` should be indented the same amount as `env` earlier in the file, and `configMapRef` should be indented the same as `configMapKeyRef`.

```
student@ckad-1:~$ vim appl/simpleapp.yaml
...
```

```
    configMapKeyRef:
      name: colors
```

```

        key: favorite
envFrom:                                #Add this and the following two lines
- configMapRef:
    name: colors
imagePullPolicy: Always
...

```

Again delete and recreate the deployment. Check that the pods restart:

```

student@ckad-1:~$ kubectl delete deployment try1
deployment.extensions "try1" deleted

student@ckad-1:~$ kubectl create -f appl/simpleapp.yaml
deployment.extensions "try1" created

```

```

student@ckad-1:~$ kubectl get pods

```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-9fnl4	1/1	Running	1	23h
registry-795c6c8b8f-hl5wf	1/1	Running	2	23h
try1-d4fbf76fd-46pkb	1/2	Running	0	40s
try1-d4fbf76fd-9kw24	1/2	Running	0	39s
try1-d4fbf76fd-bx9j9	1/2	Running	0	39s
try1-d4fbf76fd-jw8g7	1/2	Running	0	40s
try1-d4fbf76fd-lppl5	1/2	Running	0	39s
try1-d4fbf76fd-xtfd4	1/2	Running	0	40s

View the settings inside the `try1` container of a pod. The following output is truncated in a few places. Omit the container name, to observe the behavior. Also, execute a command to see all environmental variables instead of logging into the container first:

```

student@ckad-1:~$ kubectl exec -it try1-d4fbf76fd-46pkb -- /bin/bash -c
'env'
Defaulting container name to try1.
Use 'kubectl describe pod/try1-d4fbf76fd-46pkb -n default' to see all of
the containers in this pod.
REGISTRY_PORT_5000_TCP_ADDR=10.105.119.236
HOSTNAME=try1-d4fbf76fd-46pkb
TERM=xterm
yellow=y
<output_omitted>
REGISTRY_SERVICE_HOST=10.105.119.236
KUBERNETES_SERVICE_PORT=443
REGISTRY_PORT_5000_TCP=tcp://10.105.119.236:5000

```

```

KUBERNETES_SERVICE_HOST=10.96.0.1
text=black
REGISTRY_SERVICE_PORT_5000=5000
<output_omitted>
black=k
known as key

<output_omitted>
ilike=blue
<output_omitted>
magenta=m

cyan=c
<output_omitted>

```

For greater flexibility and scalability, ConfigMaps can be created from a YAML file, then deployed and redeployed as necessary. Once ingested into the cluster, the data can be retrieved in the same manner as any other object. Create another ConfigMap, this time from a YAML file:

```

student@ckad-1:~$ vim car-map.yaml
apiVersion: v1
kind: ConfigMap
metadata:
  name: fast-car
  namespace: default
data:
  car.make: Ford
  car.model: Mustang
  car.trim: Shelby

student@ckad-1:~$ kubectl create -f car-map.yaml
configmap "fast-car" created

```

View the ingested data, and note that the output is just as in file created:

```

student@ckad-1:~$ kubectl get configmap fast-car -o yaml
apiVersion: v1
data:
  car.make: Ford
  car.model: Mustang
  car.trim: Shelby
kind: ConfigMap
metadata:

```

```

creationTimestamp: 2018-04-06T16:36:32Z
name: fast-car
namespace: default
resourceVersion: "105700"
selfLink: /api/v1/namespaces/default/configmaps/fast-car
uid: aa19f8f3-39b8-11e8-ba73-42010a800003

```

Add the `configMap` settings to the `simpleapp.yaml` file as a volume. Both containers in the `try1` deployment can access the same volume, using the `volumeMounts` statements. Remember that the `volume` stanza is of equal depth to the `containers` stanza, and should probably come after for readability:

```
student@ckad-1:~$ vim appl/simpleapp.yaml
```

```

...
  spec:
    containers:
      - image: 10.105.119.236:5000/simpleapp:latest
        volumeMounts:
          - mountPath: /etc/cars
            name: car-vol
        name: car-vol
        imagePullPolicy: Always
        name: try1
...
      initialDelaySeconds: 15
      periodSeconds: 20
    Volumes:                                     #Add this and the following four lines
      - configMap:
          defaultMode: 420
          name: fast-car
        name: car-vol
        dnsPolicy: ClusterFirst
        restartPolicy: Always
...

```

Delete and recreate the deployment:

```

student@ckad-1:~$ kubectl delete deployment try1
deployment.extensions "try1" deleted

student@ckad-1:~$ kubectl create -f appl/simpleapp.yaml
deployment.extensions "try1" create

```

Verify the deployment is running. Note that we still have not automated the creation of the `/tmp/healthy` file inside the container; as a result, the **AVAILABLE** count remains zero until we use the for loop to create the file. We will remedy this in the next step.

```
student@ckad-1:~$ kubectl get deployment
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
nginx	1	1	1	1	1d
registry	1	1	1	1	1d
try1	6	6	6	0	39s

Our health check was the successful execution of a command. We will edit the command of the existing *readinessProbe* to check for the existence of the mounted `configMap` file and re-create the deployment. After a minute, both containers should become available for each pod in the deployment:

```
student@ckad-1:~$ kubectl delete deployment try1
deployment.extensions "try1" deleted
```

```
student@ckad-1:~$ vim appl/simpleapp.yaml
```

```
...
  readinessProbe:
    exec:
      command:
        - ls                                #Add this and the following line.
        - /etc/cars
      periodSeconds: 5
...

```

```
student@ckad-1:~$ kubectl create -f appl/simpleapp.yaml
deployment.extensions "try1" created
```

Wait about a minute and view the deployment and pods. All six replicas should be running and report that 2/2 containers are in a **ready** state within:

```
student@ckad-1:~$ kubectl get deployment
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
nginx	1	1	1	1	1d
registry	1	1	1	1	1d
try1	6	6	6	6	1m

```
student@ckad-1:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-9fn14	1/1	Running	1	1d

registry-795c6c8b8f-hl5wf	1/1	Running	2	1d
try1-7865dcb948-2dzc8	2/2	Running	0	1m
try1-7865dcb948-7fkh7	2/2	Running	0	1m
try1-7865dcb948-d85bc	2/2	Running	0	1m
try1-7865dcb948-djrcj	2/2	Running	0	1m
try1-7865dcb948-kwlv8	2/2	Running	0	1m
try1-7865dcb948-stb2n	2/2	Running	0	1m

View a file within the new volume mounted in a container. It should match the data we created inside the `configMap`. Because the file did not have a carriage-return, it will appear prior to the following prompt:

```
student@ckad-1:~$ kubectl exec -c try1 -it try1-7865dcb948-stb2n --
/bin/bash \
-c 'cat /etc/cars/car.trim'
Shelbystudent@ckad-1:~$
```

Attaching Storage

There are several types of storage which can be accessed with Kubernetes, with flexibility of storage being essential to scalability. In this exercise, we will configure an NFS server. With the NFS server, we will create a new *persistent volume (pv)* and a *persistent volume claim (pvc)* to use it.

Use the `CreateNFS.sh` script from the tarball to set up NFS on your master node. This script will configure the server, export `/opt/sfw` and create a file `/opt/sfw/hello.txt`.

```
student@ckad-1:~$ bash LFD259/CreateNFS.sh
Hit:1 http://us-central1.gce.archive.ubuntu.com/ubuntu xenial InRelease
Get:2 http://us-central1.gce.archive.ubuntu.com/ubuntu xenial-updates
InRelease [102 kB]
```

<output_omitted>

Should be ready. Test here and second node

```
Export list for localhost:
/opt/sfw *
```

Test by mounting the resource from your **second** node. Begin by installing the client software:

```
student@ckad-2:~$ sudo apt-get -y install nfs-common nfs-kernel-server
<output_omitted>
```

Test that you can see the exported directory using `showmount` from your second node:

```
student@ckad-2:~$ showmount -e ckad-1    ## First node's name or IP
Export list for ckad-1:
/opt/sfw *
```

Mount the directory. Be aware that, unless you edit `/etc/fstab`, this is not a persistent mount. Change out the node name for that of your master node:

```
student@ckad-2:~$ sudo mount ckad-1:/opt/sfw /mnt
```

Verify the `hello.txt` file created by the script can be viewed:

```
student@ckad-2:~$ ls -l /mnt
total 4
-rw-r--r-- 1 root root 9 Sep 28 17:55 hello.txt
```

Return to the master node and create a YAML file for an object with kind `PersistentVolume`. The included example file needs an edit to the server parameter. Use the hostname of the master server and the directory you created in the previous step. Only syntax is checked, an incorrect name or directory will not generate an error, but a Pod using the incorrect resource will not start. Note that the `accessModes` do not currently affect actual access, and are typically used as labels instead:

```
student@ckad-1:~/LFD259$ vim PVol.yaml
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pvvol-1
spec:
  capacity:
    storage: 1Gi
  accessModes:
    - ReadWriteMany
  persistentVolumeReclaimPolicy: Retain
  nfs:
    path: /opt/sfw
    server: ckad-1    #<-- Edit to match your master node name
    readOnly: false
```


Create and verify you have a new 1Gi volume named `pvvol-1`. Note the status shows as **Available**. Remember we made two persistent volumes for the image registry earlier.

```
student@ckad-1:~/LFD259$ kubectl create -f PVol.yaml
persistentvolume "pvvol-1" created
```

```
student@ckad-1:~/LFD259$ kubectl get pv
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS
CLAIM		STORAGECLASS	REASON	AGE
pvvol-1	1Gi	RWX	Retain	Available
registryvm	200Mi	RWO	Retain	Bound
default/nginx-claim0				4d
task-pv-volume	200Mi	RWO	Retain	Bound
default/registry-claim0				4d

Now that we have a new volume, we will use a persistent volume claim (pvc) to use it in a Pod. We should have two existing claims from our local registry:

```
student@ckad-1:~/LFD259$ kubectl get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES
STORAGECLASS	AGE			
nginx-claim0	Bound	registryvm	200Mi	RWO
				4d
registry-claim0	Bound	task-pv-volume	200Mi	RWO
				4d

Create a YAML file with the kind `PersistentVolumeClaim`.

```
student@ckad-1:~/LFD259$ vim pvc.yaml
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pvc-one
spec:
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 200Mi
```

Create and verify the new `pvc` status is `bound`. Note the size is 1Gi, even though 200Mi was suggested. Only a volume of at least that size could be used, so the smallest available was chosen.

```
student@ckad-1:~/LFD259$ kubectl create -f pvc.yaml
persistentvolumeclaim "pvc-one" created
```

```
student@ckad-1:~/LFD259$ kubectl get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES
STORAGECLASS	AGE			
nginx-claim0	Bound	registryvm	200Mi	RWO
4d				
pvc-one	Bound	pvvol-1	1Gi	RWX
4s				
registry-claim0	Bound	task-pv-volume	200Mi	RWO
4d				

Now, look at the status of the physical volume. It should also show as `bound`.

```
student@ckad-1:~/LFD259$ kubectl get pv
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM
STORAGECLASS	REASON	AGE			
pvvol-1	1Gi	RWX	Retain	Bound	
default/pvc-one					14m
registryvm	200Mi	RWO	Retain	Bound	
default/nginx-claim0					4d
task-pv-volume	200Mi	RWO	Retain	Bound	
default/registry-claim0					4d

Edit the `simpleapp.yaml` file to include two new sections. While one section for the container will use the volume mount point, you should have an existing entry for `car-vol`. The other section adds a volume to the deployment in general, which you can put after the `configMap` volume section.

```
student@ckad-1:~$ vim appl/simpleapp.yaml
<output_omitted>
  volumeMounts:
  - mountPath: /etc/cars
    name: car-vol
  - name: nfs-vol          ## Add these two lines
    mountPath: /opt       ##
```

```
<output_omitted>
```

```
  volumes:
```

```

- configMap:
  defaultMode: 420
  name: fast-car
  name: car-vol
- name: nfs-vol          ## Add these three lines
  persistentVolumeClaim: ##
    claimName: pvc-one   ##
<output_omitted>

```

Delete and re-create the deployment:

```

student@ckad-1:~/app1$ kubectl delete deployment try1 ; kubectl create -f \
  simpleapp.yaml
deployment.extensions "try1" deleted
deployment.extensions "try1" created

```

View the details for any of the pods in the deployment: you should see `nfs-vol` mounted under `/opt`. The use to command line completion with the tab key can be helpful for using a pod name.

```

student@ckad-1:~/app1$ kubectl describe pod try1-594fbb5fc7-5k7sj
<output_omitted>
Mounts:
  /etc/cars from car-vol (rw)
  /opt from nfs-vol (rw)
<output_omitted>

```

Rolling Updates and Rollbacks

When we started working with `simpleapp`, we used a Docker tag called `latest`. While this is the default tag when pulling an image, and commonly used, it remains just a string, and it may not be the actual latest version of the image.

Make a slight change to our source and create a new image. We will use updates and rollbacks with our application. Adding a comment to the last line should be enough for a new image to be generated:

```

student@ckad-1:~$ cd ~/app1
student@ckad-1:~/app1$ vim simple.py
<output_omitted>
## Sleep for five seconds then continue the loop

```

```
time.sleep(5)
```

```
## Adding a new comment so image is different.
```

Build the image again. A new container and image will be created. Verify when successful. There should be a different image ID and a recent creation time:

```
student@ckad-1:~/app1$ sudo docker build -t simpleapp .
Sending build context to Docker daemon 7.168 kB
Step 1/3 : FROM python:2
---> 2863c80c418c
Step 2/3 : ADD simple.py /
---> cde8ecf8492b
Removing intermediate container 3e908b76b5b4
Step 3/3 : CMD python ./simple.py
---> Running in 354620c97bf5
---> cc6bba0ea213
Removing intermediate container 354620c97bf5
Successfully built cc6bba0ea213
```

```
student@ckad-1:~/app1$ sudo docker images
```

REPOSITORY	IMAGE ID	CREATED	SIZE	TAG
simpleapp	cc6bba0ea213	8 seconds ago	679 MB	latest
10.105.119.236:5000/simpleapp	15b5ad19d313	4 days ago	679 MB	latest

<output_omitted>

Tag and push the updated image to your locally hosted registry. A reminder that your IP address will be different than the example below. Use the tag `v2` this time, instead of `latest`.

```
student@ckad-1:~/app1$ sudo docker tag simpleapp
10.105.119.236:5000/simpleapp:v2
```

```
student@ckad-1:~/app1$ sudo docker push 10.105.119.236:5000/simpleapp:v2
The push refers to a repository [10.105.119.236:5000/simpleapp]
d6153c8cc7c3: Pushed
ca82a2274c57: Layer already exists
de2fbb43bd2a: Layer already exists
4e32c2de91a6: Layer already exists
6e1b48dc2ccc: Layer already exists
ff57bdb79ac8: Layer already exists
```

```

6e5e20cbf4a7: Layer already exists
86985c679800: Layer already exists
8fad67424c4e: Layer already exists
v2: digest:
sha256:6cf74051d09463d89f1531fceb9c44cbf99006f8d9b407dd91d8f07baeee7e9c
size: 2218

```

Connect to a terminal running on your **second** node. Pull the **latest** image, then pull **v2**. Note the **latest** did not pull the new version of the image. Again, remember to use the IP for your locally hosted registry. You'll note the digest is different:

```

student@ckad-2:~$ sudo docker pull 10.105.119.236:5000/simpleapp
Using default tag: latest
latest: Pulling from simpleapp
Digest:
sha256:cefa3305c36101d32399baf0919d3482ae8a53c926688be3386f9bbc04e490a5
Status: Image is up to date for 10.105.119.236:5000/simpleapp:latest

```

```

student@ckad-2-wdrq:~$ sudo docker pull 10.105.119.236:5000/simpleapp:v2
v2: Pulling from simpleapp
f65523718fc5: Already exists
1d2dd88bf649: Already exists
c09558828658: Already exists
0e1d7c9e6c06: Already exists
c6b6fe164861: Already exists
45097146116f: Already exists
f21f8abae4c4: Already exists
1c39556edcd0: Already exists
fa67749bf47d: Pull complete
Digest:
sha256:6cf74051d09463d89f1531fceb9c44cbf99006f8d9b407dd91d8f07baeee7e9c
Status: Downloaded newer image for 10.105.119.236:5000/simpleapp:v2

```

Use **kubectl edit** to update the image for the **try1** deployment to use **v2**. As we are only changing one parameter, we could also use the **kubectl set** command. Note that the configuration file has not been updated, so a delete or a replace command would not include the new version. It can take the pods up to a minute to delete and to recreate each pod in sequence.

```

student@ckad-1:~/app1$ kubectl edit deployment try1
<output_omitted>
  containers:
  - image: 10.105.119.236:5000/simpleapp:v2
    imagePullPolicy: Always

```

<output_omitted>

Verify each of the pods has been recreated and is using the new version of the image. Note that some messages will show the scaling down of the old `replicaset`, others should show the scaling up using the new image:

```
student@ckad-1:~/app1$ kubectl get events
LAST SEEN   FIRST SEEN   COUNT   NAME
KIND        SUBOBJECT          TYPE      REASON
SOURCE      MESSAGE
4s          4s            1        try1-594fbb5fc7-nxhfx.152422073b7084da
Pod         spec.containers{goproxy} Normal    Killing
kubelet, ckad-2-wdrq Killing container with id docker://goproxy:Need to
kill Pod
<output_omitted>
2m          2m            1        try1.1524220c35a0d0fb
Deployment          Normal    ScalingReplicaSet
deployment-controller Scaled up replica set try1-895fccfb to 5
2m          2m            3        try1.1524220e0d69a94a
Deployment          Normal    ScalingReplicaSet
deployment-controller (combined from similar events): Scaled down replica
set try1-594fbb5fc7 to 0
```

View the images of a Pod in the deployment. Narrow the output to just view the images. The `goproxy` remains unchanged, but the `simpleapp` should now be `v2`:

```
student@ckad-1:~/app1$ kubectl describe pod try1-895fccfb-ttqdn |grep Image
Image:          10.105.119.236:5000/simpleapp:v2
Image ID:
docker-pullable://10.105.119.236:5000/simpleapp@sha256:6cf74051d09463d89f15
31fceb9c44cbf99006f8d9b407dd91d8f07baeee7e9c
Image:          k8s.gcr.io/goproxy:0.1
Image ID:
docker-pullable://k8s.gcr.io/goproxy@sha256:5334c7ad43048e3538775cb09aaf184
f5e8acf4b0ea60e3bc8f1d93c209865a5
```

View the update history of the deployment:

```
student@ckad-1:~/app1$ kubectl rollout history deployment try1
deployments "try1"
REVISION  CHANGE-CAUSE
```

```
1      <none>
2      <none>
```

Compare the output of the rollout history for the two revisions. Images and labels should be different, with the image `v2` being the change we made:

```
student@ckad-1:~/app1$ kubectl rollout history deployment try1 \
  --revision=1 > one.out
```

```
student@ckad-1:~/app1$ kubectl rollout history deployment try1 \
  --revision=2 > two.out
```

```
student@ckad-1:~/app1$ diff one.out two.out
1c1
< deployments "try1" with revision #1
---
> deployments "try1" with revision #2
3c3
<   Labels:      pod-template-hash=1509661973
---
>   Labels:      pod-template-hash=45197796
7c7
<     Image:      10.105.119.236:5000/simpleapp:latest
---
>     Image:      10.105.119.236:5000/simpleapp:v2
```

View what would be undone using the `--dry-run` option while undoing the rollout. This allows us to see the new template prior to using it:

```
student@ckad-1:~/app1$ kubectl rollout undo --dry-run=true deployment/try1
deployment.apps "try1"
Pod Template:
  Labels:  pod-template-hash=1509661973
          run=try1
  Containers:
    try1:
      Image: 10.105.119.236:5000/simpleapp:latest
      Port:  <none>
<output_omitted>
```

View the pods. Depending on how fast you type, the `try1` pods should be about 2 minutes old:

```
student@ckad-1:~/app1$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-9fnl4	1/1	Running	1	5d
registry-795c6c8b8f-hl5wf	1/1	Running	2	5d
try1-594fbb5fc7-7dl7c	2/2	Running	0	2m
try1-594fbb5fc7-8mxlb	2/2	Running	0	2m
try1-594fbb5fc7-jr7h7	2/2	Running	0	2m
try1-594fbb5fc7-s24wt	2/2	Running	0	2m
try1-594fbb5fc7-xfffg	2/2	Running	0	2m
try1-594fbb5fc7-zfmz8	2/2	Running	0	2m

In our case, there are only two revisions. Were there more, we could choose a particular version. The following command would have the same effect as the previous, without the `--dry-run` option.

```
student@ckad-1:~/app1$ kubectl rollout undo deployment try1 --to-revision=1
deployment.apps "try1"
```

Again, it can take a bit for the pods to be terminated and re-created. Keep checking back until they are all running again.

```
student@ckad-1:~/app1$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-9fnl4	1/1	Running	1	5d
registry-795c6c8b8f-hl5wf	1/1	Running	2	5d
try1-594fbb5fc7-7dl7c	2/2	Terminating	0	3m
try1-594fbb5fc7-8mxlb	0/2	Terminating	0	2m
try1-594fbb5fc7-jr7h7	2/2	Terminating	0	3m
try1-594fbb5fc7-s24wt	2/2	Terminating	0	2m
try1-594fbb5fc7-xfffg	2/2	Terminating	0	3m
try1-594fbb5fc7-zfmz8	1/2	Terminating	0	2m
try1-895fccfb-8dn4b	2/2	Running	0	22s
try1-895fccfb-kz72j	2/2	Running	0	10s
try1-895fccfb-rxxtw	2/2	Running	0	24s
try1-895fccfb-srwq4	1/2	Running	0	11s
try1-895fccfb-vkvmb	2/2	Running	0	31s
try1-895fccfb-z46qr	2/2	Running	0	31s

Working with Jobs

We will create a simple cron job to explore how to create them and view their execution. We will run a regular job and view both the job status and output. Note that the jobs are expected to be idempotent,

so should not be used for tasks that require strict timings to run. The `sleep 30` command will cause some jobs to finish in the next minute, as the job could start at any time during the minute.

Begin by creating a YAML file for the cron job. Set the time interval to be every minute. Use the `busybox` container and pass it the `date` command. We could just as easily use a `copy` command to back up output files from our `simpleapp`.

```
student@ckad-1:~/app1$ kubectl create -f cron-job.yaml
apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: date
spec:
  schedule: "*/1 * * * *"
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: dateperminute
              image: busybox
              args:
                - /bin/sh
                - -c
                - date; sleep 30
          restartPolicy: OnFailure
```

View the `cronjob`. Depending on the speed you type, one may not have run yet, as seen below:

```
student@ckad-1:~/app1$ kubectl get cronjob date
```

NAME	SCHEDULE	SUSPEND	ACTIVE	LAST SCHEDULE	AGE
date	*/1 * * * *	False	0	<none>	4s

View the jobs as they run. Give it a couple of minutes. Note the successful jobs completed within the timeframe of the minute, but each eventually did finish. Use `<ctrl>-c` to stop the `--watch` option.

```
student@ckad-1:~/app1$ kubectl get jobs --watch
```

NAME	DESIRED	SUCCESSFUL	AGE
date-1523426280	1	0	2s
date-1523426280	1	1	32s
date-1523426340	1	0	0s
date-1523426340	1	0	0s

`^C`

View the pods; you should see at least a couple of completed pods:

```
student@ckad-1:~/app1$ kubectl get pods | grep date
date-1523426280-zjwfm      0/1      Completed    0          1m
date-1523426340-hk897    0/1      Completed    0          42s
```

View the output of the job; you should see a recent time:

```
student@ckad-1:~/app1$ kubectl logs date-1523426340-hk897
Wed Apr 11 05:59:10 UTC 2018
```

Clean up by deleting the `cronjob`.

```
student@ckad-1:~/app1$ kubectl delete cronjob date
cronjob.batch "date" deleted
```



Lab 6.1 - Working with Security

Overview

In this lab, we will implement security features for new applications, as the `simpleapp` YAML file is getting long and more difficult to read. Kubernetes architecture favors smaller, decoupled, and transient applications, working together. We'll continue to emulate that in our exercises.

In this exercise, we will create two new applications. One will be limited in its access to the host node, but will have access to encoded data, and the second will use a network security policy to move from the default all-access Kubernetes policies to a mostly closed network. First, we will set security contexts for pods and containers, then, we will create and consume secrets, and we will finish with configuring a network security policy.

Set SecurityContext for a Pod and Container

Begin by making a new directory for our second application. Change into that directory:

```
student@ckad-1:~$ mkdir ~/app2
```

```
student@ckad-1:~$ cd ~/app2/
```

Create a YAML file for the second application. In the example below, we are using a simple image, `busybox`, which allows access to a shell, but not much more. We will add a `runAsUser` to both the pod, as well as the container:

```
student@ckad-1:~/app2$ vim second.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```

metadata:
  name: secondapp
spec:
  securityContext:
    runAsUser: 1000
  containers:
  - image: busybox
    name: secondapp
    command:
      - sleep
      - "3600"
    securityContext:
      runAsUser: 2000
      allowPrivilegeEscalation: false

```

Create the **secondapp** pod and verify it is running. Unlike the previous deployment, this application is running as a pod. Look at the YAML output, to compare and contrast with what a deployment looks like. The **status** section probably has the largest contrast:

```

student@ckad-1:~/app2$ kubectl create -f second.yaml
pod "secondapp" created

```

```

student@ckad-1:~/app2$ kubectl get pod secondapp

```

NAME	READY	STATUS	RESTARTS	AGE
secondapp	1/1	Running	0	21s

```

student@ckad-1:~/app2$ kubectl get pod secondapp -o yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: 2018-04-18T18:58:53Z
  name: secondapp
<output_omitted>

```

Execute a Bourne shell within the Pod. Check the user ID of the shell and other processes. It should show the container setting, not the pod. This allows for multiple containers within a pod to customize their UID if desired.

```

student@ckad-1:~/app2$ kubectl exec -it secondapp -- sh
/ $ ps aux

```

PID	USER	TIME	COMMAND
1	2000	0:00	sleep 3600

```

      8 2000      0:00 sh
     12 2000      0:00 ps aux

```

While here, check the capabilities of the kernel. In upcoming steps, we will modify these values.

```

/ $ grep Cap /proc/1/status
CapInh: 00000000a80425fb
CapPrm: 0000000000000000
CapEff: 0000000000000000
CapBnd: 00000000a80425fb
CapAmb: 0000000000000000
/ $ exit

```

Use the capability shell wrapper tool, the `capsh` command, to decode the output. We will view and compare the output in a few steps. Note that there are 14 comma-separated capabilities listed:

```

student@ckad-1:~/app2$ capsh --decode=00000000a80425fb
0x00000000a80425fb=cap_chown,cap_dac_override,cap_fowner,cap_fsetid,cap_kil
l,cap_setgid,cap_setuid,cap_setpcap,cap_net_bind_service,cap_net_raw,cap_sy
s_chroot,cap_mknod,cap_audit_write,cap_setfcap

```

Edit the YAML file to include new **capabilities** for the container. A capability allows granting of specific, elevated privileges without granting full root access. We will be setting **NET_ADMIN** to allow interface, routing, and other network configuration. We'll also set **SYS_TIME**, which allows system clock configuration. More on kernel capabilities can be read here:

<https://github.com/torvalds/linux/blob/master/include/uapi/linux/capability.h>

It can take up to a minute for the pod to fully terminate, allowing the future pod to be created:

```

student@ckad-1:~/app2$ kubectl delete pod secondapp
pod "secondapp" deleted

```

```

student@ckad-1:~/app2$ vim second.yaml

```

<output_omitted>

```

- sleep
- "3600"

```

securityContext:

```
runAsUser: 2000
```

```
allowPrivilegeEscalation: false
```

```
Capabilities:
```

```
add: ["NET_ADMIN", "SYS_TIME"]
```

```
Name: busy
```

#Add this and the following line

Create the pod again. Execute a shell within the container and review the `Cap` settings under `/proc/1/status`. They should be different from the previous instance:

```
student@ckad-1:~/app2$ kubectl create -f second.yaml
pod "secondapp" created

student@ckad-1:~/app2$ kubectl exec -it secondapp -- sh
/ $ grep Cap /proc/1/status
CapInh:      00000000aa0435fb
CapPrm:      0000000000000000
CapEff:      0000000000000000
CapBnd:      00000000aa0435fb
CapAmb:      0000000000000000
/ $ exit
```

Decode the output again. Note that the instance now has 16 comma-delimited capabilities listed. `cap_net_admin` is listed, as well as `cap_sys_time`.

```
student@ckad-1:~/app2$ capsh --decode=00000000aa0435fb
0x00000000aa0435fb=cap_chown,cap_dac_override,cap_fowner,cap_fsetid,cap_kill,cap_setgid,cap_setuid,cap_setpcap,cap_net_bind_service,cap_net_admin,cap_net_raw,cap_sys_chroot,cap_sys_time,cap_mknod,cap_audit_write,cap_setfcap
```

Create and Consume Secrets

Secrets are consumed in a manner similar to ConfigMaps, covered in an earlier lab. While at-rest encryption is on the way, at the moment, a secret is just base64-encoded. Begin by generating an encoded password:

```
student@ckad-1:~/app2$ echo LFTr@1n | base64
TEZUckAxbgo=
```

Create a YAML file for the object, with an API object `kind` set to `Secret`. Use the encoded key as a password parameter:

```
student@ckad-1:~/app2$ vim secret.yaml
apiVersion: v1
kind: Secret
metadata:
  name: lfsecret
```

```
data:
  password: TEZUckAxbgo=
```

Ingest the new object into the cluster:

```
student@ckad-1:~/app2$ kubectl create -f secret.yaml
secret "lfsecret" created
```

Edit the `secondapp` YAML file to use the secret as a volume mounted under `/mysqlpassword`. Note that, as there is a command executed, the pod will restart when the command finishes every 3600 seconds, or every hour.

```
student@ckad-1:~/app2$ vim second.yaml
```

<output_omitted>

```
  allowPrivilegeEscalation: false
  capabilities:
    add: ["NET_ADMIN", "SYS_TIME"]
  volumeMounts:
  - mountPath: /mysqlpassword
    name: mysql
  name: busy
  volumes:
  - name: mysql
    secret:
      secretName: lfsecret
```

```
student@ckad-1:~/app2$ kubectl delete pod secondapp
pod "secondapp" deleted
```

```
student@ckad-1:~/app2$ sleep 30 ; kubectl create -f second.yaml
pod "secondapp" created
```

Verify the pod is running, then check if the password is mounted where expected. We will find that the password is available in its clear-text, decoded state.

```
student@ckad-1:~/app2$ kubectl get pod secondapp
```

NAME	READY	STATUS	RESTARTS	AGE
secondapp	1/1	Running	0	34s

```
student@ckad-1:~/app2$ kubectl exec -ti secondapp -- /bin/sh
/ $ cat /mysqlpassword/password
LFTr@1n
```

View the location of the directory. Note that it is a symbolic link to `../data`, which is also a symbolic link to another directory. After taking a look at the filesystem within the container, exit back to the node:

```
/ $ cd /mysqlpassword/
/mysqlpassword $ ls
password
/mysqlpassword $ ls -al
total 4
drwxrwxrwt    3 root    root           100 Apr 11 07:24 .
drwxr-xr-x   21 root    root          4096 Apr 11 22:30 ..
drwxr-xr-x    2 root    root           60 Apr 11 07:24
..4984_11_04_07_24_47.831222818
lrwxrwxrwx    1 root    root           31 Apr 11 07:24 ../data ->
..4984_11_04_07_24_47.831222818
lrwxrwxrwx    1 root    root          15 Apr 11 07:24 password ->
../data/password

/mysqlpassword $ exit
```

Implement a NetworkPolicy

An early architecture decision with Kubernetes was non-isolation - that all pods were able to connect to all other pods and nodes by design. In more recent releases, the use of a **NetworkPolicy** allows for pod isolation. The policy only has effect when the network plugins, like *Project Calico*, are capable of honoring them. If used with a plugin like *flannel*, they will have no effect. The use of **matchLabels** allows for a more granular selection within the namespace, which can be selected using a **namespaceSelector**. Using multiple labels can allow for complex application of rules. More information can be found here:

<https://kubernetes.io/docs/concepts/services-networking/network-policies/>

Begin by creating a default policy which denies all traffic. Once ingested into the cluster, this will affect every pod not selected by another policy, creating a mostly-closed environment. If you want to only deny ingress or egress traffic, you can remove the other **policyType**.

```
student@ckad-1:~/app2$ vim allclosed.yaml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: deny-default
spec:
```



```
podSelector: {}
policyTypes:
- Ingress
- Egress
```

Before we can test the new network policy, we need to make sure network access works without it applied. Update `secondapp` to include a new container running `nginx`, then test access. Begin by adding two lines for the image and name `webserver`, as found below. It takes a bit for the pod to terminate, so we'll delete, then edit the file.

```
student@ckad-1:~/app2$ kubectl delete pod secondapp
pod "secondapp" deleted
```

```
student@ckad-1:~/app2$ vim second.yaml
apiVersion: v1
kind: Pod
metadata:
  name: secondapp
spec:
  securityContext:
    runAsUser: 1000
  containers:
  - image: nginx
    name: webserver
    ports:
    - containerPort: 80
  - image: busybox
    name: secondapp
    command:
<output_omitted>
```

Create the new pod. Be aware the pod will move from `ContainerCreating` to `Error` to `CrashLoopBackOff`, as only one of the containers will start. We will troubleshoot the error.

```
student@ckad-1:~/app2$ kubectl create -f second.yaml
pod "secondapp" created
```

```
student@ckad-1:~/app2$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
date-1523648520-z282m	0/1	Completed	0	3m
date-1523648580-zznjs	0/1	Completed	0	2m
date-1523648640-2h6qt	0/1	Completed	0	1m
date-1523648700-drfp5	1/1	Running	0	18s

nginx-6b58d9cdfd-9fnl4	1/1	Running	1	8d
registry-795c6c8b8f-hl5wf	1/1	Running	2	8d
secondapp	1/2	CrashLoopBackOff	1	13s

Take a closer look at the events leading up to the failure. The images were pulled, but it was the execution of the container which failed.

```
student@ckad-1:~/app2$ kubectl get event
<output_omitted>
      Normal          Created                  kubelet, ckad-2-wdrq      Created
container
5m          5m          1          secondapp.1525166dae0e43
Pod          spec.containers{busy}          Normal      Started
kubelet, ckad-2-wdrq      Started container
20s          5m          25          secondapp.1525166e5791a7fd
Pod          spec.containers{webserver}          Warning    BackOff
kubelet, ckad-2-wdrq      Back-off restarting failed container
```

View the logs of the **webserver** container mentioned in the previous output. Note that there are errors about the user directive and not having permission to make directories.

```
student@ckad-1:~/app2$ kubectl logs secondapp webserver
2018/04/13 19:51:13 [warn] 1#1: the "user" directive makes sense only if
the master process runs with super-user privileges, ignored in
/etc/nginx/nginx.conf:2
nginx: [warn] the "user" directive makes sense only if the master process
runs with super-user privileges, ignored in /etc/nginx/nginx.conf:2
2018/04/13 19:51:13 [emerg] 1#1: mkdir() "/var/cache/nginx/client_temp"
failed (13: Permission denied)
nginx: [emerg] mkdir() "/var/cache/nginx/client_temp" failed (13:
Permission denied)
```

Delete the pods. Edit the YAML file to comment out the setting of a UID for the entire pod.

```
student@ckad-1:~/app2$ kubectl delete -f second.yaml
pod "secondapp" deleted
```

```
student@ckad-1:~/app2$ vim second.yaml
<output_omitted>
  name: secondapp
spec:
#   securityContext:
#     runAsUser: 1000
```

```

containers:
- image: nginx
  name: webserver
  ports:
  - containerPort: 80
<output_omitted>

```

Create the pod again. This time, both containers should run. You may have to wait for the previous pod to fully terminate, depending on how fast you type.

```

student@ckad-1:~/app2$ kubectl create -f second.yaml
pod "secondapp" created

```

```

student@ckad-1:~/app2$ kubectl get pods
NAME                                READY    STATUS    RESTARTS   AGE
secondapp                          2/2      Running    0           5s

```

Expose the **webserver** using a **NodePort** service. Expect an error due to the lack of labels:

```

student@ckad-1:~/app2$ kubectl expose pod secondapp --type=NodePort
--port=80
error: couldn't retrieve selectors via --selector flag or introspection:
the pod has no labels and cannot be exposed
See 'kubectl expose -h' for help and examples.

```

Edit the YAML file to add a label in the metadata, adding the **example: second** label right after the pod name. Note that you can delete several resources at once by passing the YAML file to the **delete** command. Delete and recreate the pod. It may take up to a minute for the pod to shut down:

```

student@ckad-1:~/app2$ kubectl delete -f second.yaml
pod "secondapp" deleted

```

```

student@ckad-1:~/app2$ vim second.yaml
apiVersion: v1
kind: Pod
metadata:
  name: secondapp
  labels:
    example: second
spec:
#  securityContext:
#    runAsUser: 1000

```

<output_omitted>

```
student@ckad-1:~/app2$ kubectl create -f second.yaml
pod "secondapp" created
```

```
student@ckad-1:~/app2$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
<output_omitted>				
secondapp	2/2	Running	0	15s

This time we will expose the `NodePort` again, and create the service separately, then add a label to illustrate how labels are essential for tying resources together inside of Kubernetes.

```
student@ckad-1:~/app2$ kubectl create service nodeport secondapp --tcp=80
service "secondapp" created
```

Look at the details of the service. Note the selector is set to `app: secondapp`. Also, take the note of the `nodePort`, which is 31655 in the example below; yours may be different:

```
student@ckad-1:~/app2$ kubectl get svc secondapp -o yaml
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: 2018-04-19T22:07:25Z
  labels:
    app: secondapp
  name: secondapp
  namespace: default
  resourceVersion: "216490"
  selfLink: /api/v1/namespaces/default/services/secondapp
  uid: 0aeaea82-441e-11e8-ac6e-42010a800007
spec:
  clusterIP: 10.97.96.75
  externalTrafficPolicy: Cluster
  ports:
  - name: "80"
    nodePort: 31655
    port: 80
    protocol: TCP
    targetPort: 80
  selector:
    app: secondapp
  sessionAffinity: None
```

```

    type: NodePort
status:
  loadBalancer: {}

```

Test access to the service. As the label does not match any other resources, the `curl` command should hang, and eventually time out.

```
student@ckad-1:~/app2$ curl http://10.97.96.75
```

Edit the service. We will change the label to match `secondapp`, and set the `nodePort` to a new port, one that may have been specifically opened by our firewall team, port 32000.

```

student@ckad-1:~/app2$ kubectl edit svc secondapp
<output_omitted>
ports:
- name: "80"
  nodePort: 32000      ## Edit this line
  port: 80
  protocol: TCP
  targetPort: 80
selector:
  example: second      ## Edit this line, too
  sessionAffinity: None
<output_omitted>

```

Verify the updated port number is showing properly, and take note of the `ClusterIP`. The example below shows a `clusterIP` of 10.97.96.75 and a port of 32000 as expected.

```

student@ckad-1:~/app2$ kubectl get svc
NAME                TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
<output_omitted>
secondapp           NodePort    10.97.96.75    <none>         80:32000/TCP     5m

```

Test access to the high port. You should get the default `nginx` page both if you test from the node to the `ClusterIP:low-port` and from the exterior `hostIP:highport`. As the high port is randomly generated, make sure it is available:

```

student@ckad-1:~/app2$ curl http://10.97.96.75
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>

```

```
[serewicz@laptop ~]$ curl http://35.184.219.5:32000
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>
```

Now, test egress from a container to the outside world. We'll use the `netcat` command to verify access to a running webserver on port 80. First, test local access to nginx, then a remote server.

```
student@ckad-1:~/app2$ kubectl exec -it -c busy secondapp sh
/ $ nc -vz 127.0.0.1 80
127.0.0.1 (127.0.0.1:80) open

/ $ nc -vz www.linux.com 80
www.linux.com (151.101.185.5:80) open

/ $ exit
```

Now that we have tested both ingress and egress, we can implement the network policy:

```
student@ckad-1:~/app2$ kubectl create -f allclosed.yaml
networkpolicy.networking.k8s.io "deny-default" created
```

Use the ingress and egress tests again. Three of the four should eventually time out. Start by testing from outside the cluster:

```
[serewicz@laptop ~]$ curl http://35.184.219.5:32215
curl: (7) Failed to connect to 35.184.219.5 port 32000: Connection timed out
```

Then, test from the host to the container:

```
student@ckad-1:~/app2$ curl http://10.97.96.75:80
curl: (7) Failed to connect to 10.97.96.75 port 80: Connection timed out
```

Now, test egress. From container to container should work, as the filter is outside of the pod; then, test egress. It should time out:

```
student@ckad-1:~/app2$ kubectl exec -it -c busy secondapp sh
/ $ nc -vz 127.0.0.1 80
127.0.0.1 (127.0.0.1:80) open
```

```

/ $ nc -vz www.linux.com 80
nc: bad address 'www.linux.com'

/ $ exit

```

Update the `NetworkPolicy` and remove the `Egress` line. Then, replace the policy:

```

student@ckad-1:~/app2$ vim allclosed.yaml

student@ckad-1:~/app2$ kubectl replace -f allclosed.yaml
networkpolicy.networking.k8s.io "deny-default" replaced

```

Test egress access to an outside site. Get the IP address of the `eth0` inside the container while logged in. The IP is `192.168.55.91` in the example below, but yours may be different:

```

student@ckad-1:~/app2$ kubectl exec -it -c busy secondapp sh
/ $ nc -vz www.linux.com 80
www.linux.com (151.101.185.5:80) open
/ $ ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: tunl0@NONE: <NOARP> mtu 1480 qdisc noop qlen 1000
    link/ipip 0.0.0.0 brd 0.0.0.0
4: eth0@if59: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1500 qdisc
noqueue
    link/ether 1e:c8:7d:6a:96:c3 brd ff:ff:ff:ff:ff:ff
    inet 192.168.55.91/32 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::1cc8:7dff:fe6a:96c3/64 scope link
        valid_lft forever preferred_lft forever
/ $ exit

```

Now, add a selector to allow ingress to only the `nginx` container. Use the IP from the `eth0` range:

```

student@ckad-1:~/app2$ vim allclosed.yaml
<output_omitted>
policyTypes:
- Ingress

```

```

ingress:
- from:
  - ipBlock:
      cidr: 192.168.0.0/16

```

Recreate the policy, and verify its configuration.

```

student@ckad-1:~/app2$ kubectl replace -f allclosed.yaml
networkpolicy.networking.k8s.io "deny-default" replaced

```

```

student@ckad-1:~/app2$ kubectl get networkpolicy
NAME                POD-SELECTOR      AGE
deny-default        example=second    15s

```

```

student@ckad-1:~/app2$ kubectl get networkpolicy -o yaml
apiVersion: v1
items:
- apiVersion: extensions/v1beta1
  kind: NetworkPolicy
  metadata:
<output_omitted>

```

Test access to the container both using `curl`, as well as `ping`, the IP address to use was found from `ip a` inside the container:

```

student@ckad-1:~/app2$ curl http://192.168.55.91
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>

```

```

student@ckad-1:~/app2$ ping -c5 192.168.55.91
PING 192.168.55.91 (192.168.55.91) 56(84) bytes of data.
64 bytes from 192.168.55.91: icmp_seq=1 ttl=63 time=1.11 ms
64 bytes from 192.168.55.91: icmp_seq=2 ttl=63 time=0.352 ms
64 bytes from 192.168.55.91: icmp_seq=3 ttl=63 time=0.350 ms
64 bytes from 192.168.55.91: icmp_seq=4 ttl=63 time=0.359 ms
64 bytes from 192.168.55.91: icmp_seq=5 ttl=63 time=0.295 ms

--- 192.168.55.91 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4054ms

```

```
rtt min/avg/max/mdev = 0.295/0.495/1.119/0.312 ms
```

Update the policy to only allow ingress for TCP traffic on port 80, then test with `curl`, which should work:

```
student@ckad-1:~/app2$ vim allclosed.yaml
```

```
<output_omitted>
```

```
- Ingress
```

```
  ingress:
```

```
- from:
```

```
  - ipBlock:
```

```
    cidr: 192.168.0.0/16
```

```
  ports:
```

```
- port: 80
```

```
  protocol: TCP
```

```
student@ckad-1:~/app2$ kubectl replace -f allclosed.yaml
networkpolicy.networking.k8s.io "deny-default" replaced
```

```
student@ckad-1:~/app2$ curl http://192.168.55.91
```

```
<!DOCTYPE html>
```

```
<html>
```

```
<head>
```

```
<title>Welcome to nginx!</title>
```

```
<output_omitted>
```

All five pings should fail, with zero received:

```
student@ckad-1:~/app2$ ping -c5 192.168.55.91
```

```
PING 192.168.55.91 (192.168.55.91) 56(84) bytes of data.
```

```
--- 192.168.55.91 ping statistics ---
```

```
5 packets transmitted, 0 received, 100% packet loss, time 4098ms
```



Lab 7.1 - Exposing Applications

Overview

In this lab, we will explore various ways to expose an application to other pods and outside the cluster. We will add to the **NodePort** used in previous labs and other service options.

Expose A Service

We will begin by using the default service type **ClusterIP**. This is a cluster internal IP, only reachable from within the cluster. Begin by viewing the existing services:

```
student@ckad-1:~$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	8d
nginx	ClusterIP	10.108.95.67	<none>	443/TCP	8d
registry	ClusterIP	10.105.119.236	<none>	5000/TCP	8d
secondapp	NodePort	10.111.26.8	<none>	80:32000/TCP	7h

Delete the existing service for **secondapp**.

```
student@ckad-1:~/app2$ kubectl delete svc secondapp
service "secondapp" deleted
```

Create a YAML file for a replacement service, which would be persistent. Use the label to select the **secondapp**. Expose the same port and protocol of the previous service:

```
student@ckad-1:~/app2$ vim service.yaml
apiVersion: v1
```

```
kind: Service
metadata:
  name: secondapp
  labels:
    run: my-nginx
spec:
  ports:
  - port: 80
    protocol: TCP
  selector:
    example: second
```

Create the service, find the new IP and port. Note that there is no high number port, as this is internal access only:

```
student@ckad-1:~/app2$ kubectl create -f service.yaml
service "secondapp" created
```

```
student@ckad-1:~/app2$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	8d
nginx	ClusterIP	10.108.95.67	<none>	443/TCP	8d
registry	ClusterIP	10.105.119.236	<none>	5000/TCP	8d
secondapp	ClusterIP	10.98.148.52	<none>	80/TCP	14s

Test access. You should see the default welcome page again:

```
student@ckad-1:~/app2$ curl http://10.98.148.52
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>
```

To expose a port to outside the cluster, we will create a **NodePort**. We had done this in a previous step from the command line. When we create a **NodePort**, it will create a new **ClusterIP** automatically. Edit the YAML file again. Add **type: NodePort**. Also, add the high-port to match an open port in the firewall, as mentioned in the previous chapter. You'll have to delete and re-create, as the existing IP is immutable, but not able to be reused. The **NodePort** will try to create a new **ClusterIP** instead.

```
student@ckad-1:~/app2$ vim service.yaml
apiVersion: v1
```

```
kind: Service
metadata:
  name: secondapp
  labels:
    run: my-nginx
spec:
  ports:
  - port: 80
    protocol: TCP
    nodePort: 32000
  type: NodePort
  selector:
    example: second
```

```
student@ckad-1:~/app2$ kubectl delete svc secondapp ; kubectl create \
-f service.yaml
service "secondapp" deleted
service "secondapp" created
```

Find the new `ClusterIP` and ports for the service:

```
student@ckad-1:~/app2$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	8d
nginx	ClusterIP	10.108.95.67	<none>	443/TCP	8d
registry	ClusterIP	10.105.119.236	<none>	5000/TCP	8d
secondapp	NodePort	10.109.134.221	<none>	80:32000/TCP	4s

Test the low port number using the `ClusterIP`:

```
student@ckad-1:~/app2$ curl 10.109.134.221
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>
```

Test access from an external node to the host IP and the high container port. Your IP and port will be different. It should work, even with the network policy in place, as the traffic is arriving via a 192.168.0.0 port.

```
serewicz@laptop:~/Desktop$ curl http://35.184.219.5:32000
<!DOCTYPE html>
```

```
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>
```

The use of a **LoadBalancer** makes an asynchronous request to an external provider for a load balancer, if one is available. It then creates a **NodePort** and waits for a response, including the external IP. The local **NodePort** will work even before the load balancer replies. Edit the YAML file and change the **type** to be **LoadBalancer**.

```
student@ckad-1:~/app2$ vim service.yaml
```

```
<output_omitted>
```

```
- port: 80
  protocol: TCP
  type: LoadBalancer
  selector:
    example: second
```

```
student@ckad-1:~/app2$ kubectl delete svc secondapp
service "secondapp" deleted
```

```
student@ckad-1:~/app2$ kubectl create -f service.yaml
service "secondapp" created
```

As mentioned, the cloud provider is not configured to provide a load balancer; the **External-IP** will remain in pending state. Some issues have been found using this with VirtualBox.

```
student@ckad-1:~/app2$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
AGE				
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP
8d				
nginx	ClusterIP	10.108.95.67	<none>	443/TCP
8d				
registry	ClusterIP	10.105.119.236	<none>	5000/TCP
8d				
secondapp	LoadBalancer	10.109.26.21	<pending>	80:32000/TCP
4s				

Test again local and from a remote node. The IP addresses and ports will be different on your node.

```
serewic@laptop:~/Desktop$ curl http://35.184.219.5:32000
```

```
<!DOCTYPE html>
```

```
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>
```

Ingress Controller

If you have a large number of services to expose outside of the cluster, or to expose a low-number port on the host node, you can deploy an ingress controller. While nginx and GCE have controllers officially supported by Kubernetes.io, the Traefik Ingress Controller is easier to install, at least at the moment.

As we have RBAC configured, we need to make sure the controller will run and be able to work with all necessary ports, endpoints, and resources. Create a YAML file to declare a `clusterrole` and a `clusterrolebinding`:

```
student@ckad-1:~/app2$ vim ingress.rbac.yaml
kind: ClusterRole
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: traefik-ingress-controller
rules:
- apiGroups:
  - ""
  resources:
  - services
  - endpoints
  - secrets
  verbs:
  - get
  - list
  - watch
- apiGroups:
  - extensions
  resources:
  - ingresses
  verbs:
  - get
  - list
  - watch
```

```

---
kind: ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: traefik-ingress-controller
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: traefik-ingress-controller
subjects:
- kind: ServiceAccount
  name: traefik-ingress-controller
  namespace: kube-system

```

Create the new role and binding:

```

student@ckad-1:~/app2$ kubectl create -f ingress.rbac.yaml
clusterrole.rbac.authorization.k8s.io "traefik-ingress-controller" created
clusterrolebinding.rbac.authorization.k8s.io "traefik-ingress-controller"
created

```

Create the Traefik controller. We will use a script directly from their website:

```

student@ckad-1:~/app2$ kubectl apply -f
https://raw.githubusercontent.com/containous/traefik/master/examples/k8s/tr
aefik-ds.yaml
serviceaccount "traefik-ingress-controller" created
daemonset.extensions "traefik-ingress-controller" created
service "traefik-ingress-service" created

```

Now that there is a new controller, we need to pass some rules, so it knows how to handle requests. Note that the host mentioned is www.example.com, which is probably not your node name. We will pass a false header when testing. Also, the service name needs to match the **secondapp** we've been working with:

```

student@ckad-1:~/app2$ vim ingress.rule.yaml
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: ingress-test
  annotations:
    kubernetes.io/ingress.class: traefik
spec:

```

```
rules:
- host: www.example.com
  http:
    paths:
    - backend:
        serviceName: secondapp
        servicePort: 80
      path: /
```

Now, ingest the rule into the cluster:

```
student@ckad-1:~/app2$ kubectl create -f ingress.rule.yaml
ingress.extensions "ingress-test" created
```

We should be able to test the internal and external IP addresses, and see the nginx welcome page. The loadbalancer would present the traffic, a `curl` request in this case, to the externally facing interface. Use `ip a` to find the IP address of the interface which would face the loadbalancer. In this example, the interface would be `ens4`, and the IP would be 10.128.0.7.

```
student@ckad-1:~$ ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group
default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: ens4: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1460 qdisc mq state UP group
default qlen 1000
    link/ether 42:01:0a:80:00:03 brd ff:ff:ff:ff:ff:ff
    inet 10.128.0.3/32 brd 10.128.0.3 scope global ens4
        valid_lft forever preferred_lft forever
<output_omitted>
```

```
student@ckad-1:~/app2$ curl -H "Host: www.example.com" http://10.128.0.7/
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
```

```
serewicz@laptop:~$ curl -H "Host: www.example.com" http://35.193.3.179
<!DOCTYPE html>
```



```
<html>
<head>
<title>Welcome to nginx!</title>
<style>
<output_omitted>
```

At this point, we would keep adding more and more web servers. We'll configure one more, which would then be a process continued as many times as desired.

Begin by deploying another `nginx` server. Give it a label and expose port 80:

```
student@ckad-1:~/app2$ kubectl run thirdpage --image=nginx \
  --port=80 -l example=third
deployment.apps "thirdpage" created
```

Expose the new server as a `NodePort`:

```
student@ckad-1:~/app2$ kubectl expose deployment thirdpage --type=NodePort
service "thirdpage" exposed
```

Now, we will customize the installation. Run a bash shell inside the new pod. Your pod name will end differently. Install `vim` inside the container, then edit the `index.html` file of `nginx` so that the title of the web page will be *Third Page*.

```
student@ckad-1:~/app2$ kubectl exec -it thirdpage-5cf8d67664-zcmfh --
/bin/bash
```

```
root@thirdpage-5cf8d67664-zcmfh:/# apt-get update
```

```
root@thirdpage-5cf8d67664-zcmfh:/# apt-get install vim -y
```

```
root@thirdpage-5cf8d67664-zcmfh:/# vim /usr/share/nginx/html/index.html
<!DOCTYPE html>
<html>
<head>
<title>Third Page</title>
<style>
<output_omitted>
```

Edit the ingress rules to point the `thirdpage` service:

```
student@ckad-1:~/app2$ kubectl edit ingress ingress-test
<output_omitted>
- host: www.example.com
  http:
    paths:
      - backend:
          serviceName: secondapp
          servicePort: 80
        path: /
- host: thirdpage.org
  http:
    paths:
      - backend:
          serviceName: thirdpage
          servicePort: 80
        path: /
status:
<output_omitted>
```

Test the second hostname using `curl` locally, as well as from a remote system:

```
student@ckad-1:~/app2$ curl -H "Host: thirdpage.org" http://10.128.0.7/
<!DOCTYPE html>
<html>
<head>
<title>Third Page</title>
<style>
<output_omitted>
```



Lab 8.1 - Troubleshooting

Monitor Applications

View the **secondapp** pod, it should show as running. This may not mean the application within is working properly, but that the pod is running. The restarts are due to the command we have written to run. The pod exists when done, and the controller restarts another container inside. The count depends on how long the labs have been running.

```
student@ckad-1/app1:~$ cd
student@ckad-1:~$ kubectl get pods secondapp
```

NAME	READY	STATUS	RESTARTS	AGE
secondapp	2/2	Running	49	2d

Look closer at the pod. Working slowly through the output, check each line. If you have issues, are other pods having issues on the same node or volume? Check the state of each container. Both **busy** and **webserver** should report **Running**. Note that **webserver** has a restart count of zero, while **busy** has a restart count of 49. We expect this, as in our case, the pod has been running for 49 hours.

```
student@ckad-1:~$ kubectl describe pod secondapp
```

Name: secondapp
 Namespace: default
 Node: ckad-2-wdrq/10.128.0.2
 Start Time: Fri, 13 Apr 2018 20:34:56 +0000
 Labels: example=second
 Annotations: <none>
 Status: Running
 IP: 192.168.55.91
 Containers:
 webserver:
 <output_omitted>

```

State:           Running
  Started:       Fri, 13 Apr 2018 20:34:58 +0000
  Ready:         True
  Restart Count: 0
<output_omitted>

```

```

busy:
<output_omitted>

```

```

State:           Running
  Started:       Sun, 15 Apr 2018 21:36:20 +0000
  Last State:    Terminated
    Reason:      Completed
    Exit Code:   0
  Started:       Sun, 15 Apr 2018 20:36:18 +0000
  Finished:      Sun, 15 Apr 2018 21:36:18 +0000
  Ready:         True
  Restart Count: 49
  Environment:   <none>

```

There are three values for conditions. Check that the pod reports **Initialized**, **Ready** and **Scheduled**.

```

<output_omitted>
Conditions:
  Type           Status
  Initialized     True
  Ready          True
  PodScheduled   True
<output_omitted>

```

Check if there are any events with errors or warnings which may indicate what is causing any problems:

```

Events:
  Type    Reason   Age           From          Message
  ----    -
  Normal  Pulling   34m (x50 over 2d)  kubelet, ckad-2-wdrq  pulling image "busybox"
  Normal  Pulled    34m (x50 over 2d)  kubelet, ckad-2-wdrq  Successfully pulled image "busybox"
  Normal  Created   34m (x50 over 2d)  kubelet, ckad-2-wdrq  Created container

```

Normal Started 34m (x50 over 2d) kubelet, ckad-2-wdrq Started container

View each container log. You may have to sift errors from expected output. Some containers may have no output at all, as is found with **busy**.

```
student@ckad-1:~$ kubectl logs secondapp webserver
192.168.55.0 - - [13/Apr/2018:21:18:13 +0000] "GET / HTTP/1.1" 200 612 "-"
"curl/7.47.0" "-"
192.168.55.0 - - [13/Apr/2018:21:20:35 +0000] "GET / HTTP/1.1" 200 612 "-"
"curl/7.53.1" "-"
127.0.0.1 - - [13/Apr/2018:21:25:29 +0000] "GET" 400 174 "-" "-" "-"
127.0.0.1 - - [13/Apr/2018:21:26:19 +0000] "GET index.html" 400 174 "-" "-"
"_"
<output_omitted>
```

```
student@ckad-1:~$ kubectl logs secondapp busy
student@ckad-1:~$
```

Check to make sure the container is able to use DNS and communicate with the outside world. Remember that we still have limited the UID for **secondapp** to be UID 2000, which may prevent some commands from running. It can also prevent an application from completing expected tasks:

```
student@ckad-1:~$ kubectl exec -it secondapp -c busy -- sh
/ $ nslookup www.linux.com
Server:      10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local

Name:        www.linux.com
Address 1: 151.101.45.5

/ $ cat /etc/resolv.conf
nameserver 10.96.0.10
search default.svc.cluster.local svc.cluster.local cluster.local
c.endless-station-188822.internal google.internal
options ndots:5
```

Test access to a remote node using **NetCat**. There are several options to **nc** which can help troubleshoot if the problem is the local node, something between systems or the target. In the example below, the connect never completes and a **<ctrl>-c** was used to interrupt.

```
/ $ nc www.linux.com 25
```

^Cpunt!

Test using an IP address in order to narrow the issue to name resolution. In this case, the IP in use is a well known IP for Google's DNS servers. The following example shows that Internet name resolution is working, but our UID issue prevents access to the `index.html` file.

```
/ $ wget http://www.linux.com/
Connecting to www.linux.com (151.101.45.5:80)
Connecting to www.linux.com (151.101.45.5:443)
wget: can't open 'index.html': Permission denied
```

```
/ $ exit
```

Make sure traffic is being sent to the correct Pod. Check the details of both the service and endpoint. Pay close attention to ports in use, as a simple typo can prevent traffic from reaching the proper pod. Make sure labels and selectors don't have any typos as well.

```
student@ckad-1:~$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
AGE				
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP
10d				
nginx	ClusterIP	10.108.95.67	<none>	443/TCP
10d				
registry	ClusterIP	10.105.119.236	<none>	5000/TCP
10d				
secondapp	LoadBalancer	10.109.26.21	<pending>	80:32000/TCP
1d				
thirdpage	NodePort	10.109.250.78	<none>	80:31230/TCP
1h				

```
student@ckad-1:~$ kubectl get svc secondapp -o yaml
```

```
<output_omitted>
```

```
clusterIP: 10.109.26.21
externalTrafficPolicy: Cluster
ports:
- nodePort: 32000
  port: 80
  protocol: TCP
  targetPort: 80
selector:
  example: second
```

<output_omitted>

Verify an endpoint for the service exists and has expected values, including namespaces, ports and protocols.

```
student@ckad-1:~$ kubectl get ep
NAME                ENDPOINTS                AGE
kubernetes          10.128.0.3:6443          10d
nginx                192.168.55.68:443        10d
registry            192.168.55.69:5000       10d
secondapp            192.168.55.91:80         1d
thirdpage            192.168.241.57:80        1h

student@ckad-1:~$ kubectl get ep secondapp -o yaml
apiVersion: v1
kind: Endpoints
metadata:
  creationTimestamp: 2018-04-14T05:37:32Z
<output_omitted>
```

If the containers, services and endpoints are working, the issue may be with an infrastructure service like `kube-proxy`. Ensure it's running, then look for errors in the logs. As we have two nodes, we will have two proxies to look at. As we built our cluster with `kubeadm`, the proxy runs as a container. On other systems, you may need to use `journalctl` or look under `/var/log/kube-proxy.log`.

```
student@ckad-1:~$ ps -elf |grep kube-proxy
4 S root      2864  2847  0  80   0 - 14178 -      15:45 ?          00:00:56
/usr/local/bin/kube-proxy --config=/var/lib/kube-proxy/config.conf
0 S student  23513 18282  0  80   0 - 3236 pipe_w 22:49 pts/0    00:00:00
grep --color=auto kube-proxy
```

```
student@ckad-1:~$ journalctl -a | grep proxy
Apr 15 15:44:43 ckad-2-nzjr audit[742]: AVC apparmor="STATUS"
operation="profile_load" profile="unconfined"
name="/usr/lib/lxd/lxd-bridge-proxy" pid=742 comm="apparmor_parser"
Apr 15 15:44:43 ckad-2-nzjr kernel: audit: type=1400
audit(1523807083.011:11): apparmor="STATUS" operation="profile_load"
profile="unconfined" name="/usr/lib/lxd/lxd-bridge-proxy" pid=742
comm="apparmor_parser"
Apr 15 15:45:17 ckad-2-nzjr kubelet[1248]: I0415 15:45:17.153670      1248
reconciler.go:217] operationExecutor.VerifyControllerAttachedVolume started
for volume "xtables-lock" (UniqueName:
```

```
"kubernetes.io/host-path/e701fc01-38f3-11e8-a142-42010a800003-xtables-lock"
) pod "kube-proxy-t8k4w" (UID: "e701fc01-38f3-11e8-a142-42010a800003")
```

Look at both of the proxy logs. Lines which begin with the character `I` are info, `E` are errors. In this example, the last message says access to listing an endpoint was denied by RBAC. It was because a default installation via Helm wasn't RBAC-aware. If you are not using the command line completion, view the possible pod names first.

```
student@ckad-1:~$ kubectl -n kube-system get pod
```

```
student@ckad-1:~$ kubectl -n kube-system logs kube-proxy-fsdf
I0405 17:28:37.091224      1 feature_gate.go:190] feature gates: map[]
W0405 17:28:37.100565      1 server_others.go:289] Flag proxy-mode=""
unknown, assuming iptables proxy
I0405 17:28:37.101846      1 server_others.go:138] Using iptables Proxier.
I0405 17:28:37.121601      1 server_others.go:171] Tearing down inactive
rules.
<output_omitted>
E0415 15:45:17.086081      1 reflector.go:205]
k8s.io/kubernetes/pkg/client/informers/informers_generated/internalversion/
factory.go:85: Failed to list *core.Endpoints: endpoints is forbidden: User
"system:serviceaccount:kube-system:kube-proxy" cannot list endpoints at the
cluster scope: [clusterrole.rbac.authorization.k8s.io "system:node-proxier"
not found, clusterrole.rbac.authorization.k8s.io "system:basic-user" not
found, clusterrole.rbac.authorization.k8s.io "system:discovery" not found]
```

Check that the proxy is creating the expected rules for the problem service. Find the destination port being used for the service, 30195 in this case.

```
student@ckad-1:~$ sudo iptables-save |grep secondapp
-A KUBE-NODEPORTS -p tcp -m comment --comment "default/secondapp:" -m tcp
--dport 30195 -j KUBE-MARK-MASQ
-A KUBE-NODEPORTS -p tcp -m comment --comment "default/secondapp:" -m tcp
--dport 30195 -j KUBE-SVC-DAASHM5XQZF5XI3E
-A KUBE-SEP-YDKKGXN54FN2TFPE -s 192.168.55.91/32 -m comment --comment
"default/secondapp:" -j KUBE-MARK-MASQ
-A KUBE-SEP-YDKKGXN54FN2TFPE -p tcp -m comment --comment
"default/secondapp:" -m tcp -j DNAT --to-destination 192.168.55.91:80
-A KUBE-SERVICES ! -s 192.168.0.0/16 -d 10.109.26.21/32 -p tcp -m comment
--comment "default/secondapp: cluster IP" -m tcp --dport 80 -j
KUBE-MARK-MASQ
```



```
-A KUBE-SERVICES -d 10.109.26.21/32 -p tcp -m comment --comment
"default/secondapp: cluster IP" -m tcp --dport 80 -j
KUBE-SVC-DAASHM5XQZF5XI3E
-A KUBE-SVC-DAASHM5XQZF5XI3E -m comment --comment "default/secondapp:" -j
KUBE-SEP-YDKKGXN54FN2TFPE
-A KUBE-SVC-NPX46M4PTMTKRN6Y -m comment --comment
"default/kubernetes:https" -m recent --rcheck --seconds 10800 --reap --name
KUBE-SEP-2QXHNT77UCWCSQLV --mask 255.255.255.255 --rsource -j
KUBE-SEP-2QXHNT77UCWCSQLV
```

Ensure the proxy is working by checking the port targeted by iptables. If it fails, open a second terminal and view the proxy logs when making a request as it happens.

```
student@ckad-1:~$ curl localhost:32000
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>
```

Conformance Testing

The CNCF group is in the process of formalizing what is considered to be a conforming Kubernetes cluster. While that project matures, there is an existing tool provided by Heptio which can be useful. We will need to make sure a newer version of Golang is installed for it to work. You can download the code from GitHub and look around with git or with go, depending on which tool you are most familiar.

Create a new directory to hold the testing code:

```
student@ckad-1:~$ mkdir test

student@ckad-1:~$ cd test/
```

Use `git` to download the Sonobuoy code. View the resource after it downloads:

```
student@ckad-1:~/test$ git clone https://github.com/heptio/sonobuoy
Cloning into 'sonobuoy'...
remote: Counting objects: 7847, done.
remote: Compressing objects: 100% (23/23), done.
remote: Total 7847 (delta 2), reused 0 (delta 0), pack-reused 7824
```

```
Receiving objects: 100% (7847/7847), 10.19 MiB | 0 bytes/s, done.
Resolving deltas: 100% (3818/3818), done.
Checking connectivity... done.
```

```
student@ckad-1:~/test$ ls
sonobuoy
```

```
student@ckad-1:~/test$ cd sonobuoy/
```

```
student@ckad-1:~/test/sonobuoy$ ls
cmd                Gopkg.toml          scripts
CODE_OF_CONDUCT.md heptio-images-ee4b0474b93e.json.enc SUPPORT.md
CONTRIBUTING.md    LICENSE             test
Dockerfile          main.go             travis-deploy.sh
docs                Makefile            vendor
examples            pkg
Gopkg.lock          README.md
```

```
student@ckad-1:~/test/sonobuoy$ less README.md
```

The Heptio team suggests the use of an easy-to-use Golang tool `gimme`. We will follow their suggestion and use it to pull their code. Start by making sure you have a `/bin` directory under your home directory.

```
student@ckad-1:~/test/sonobuoy$ mkdir ~/bin
```

Use `curl` to download the binary. Note the use of `-o` as in output to save the binary to the newly created directory:

```
student@ckad-1:~/test/sonobuoy$ curl -sL -o ~/bin/gimme \
https://raw.githubusercontent.com/travis-ci/gimme/master/gimme
```

View the file. Note it is not yet executable. Make it so:

```
student@ckad-1:~/test/sonobuoy$ ls -l ~/bin/gimme
-rw-rw-r-- 1 student student 27035 Apr 15 20:46 /home/student/bin/gimme
```

```
student@ckad-1:~/test/sonobuoy$ chmod +x ~/bin/gimme
```

Use the `gimme` tool to download the stable version of Go:

```
student@ckad-1:~/test/sonobuoy$ ~/bin/gimme stable
```

```
unset GOOS;
unset GOARCH;
export GOROOT='/home/student/.gimme/versions/go1.10.1.linux.amd64';
export
PATH="/home/student/.gimme/versions/go1.10.1.linux.amd64/bin:${PATH}";
go version >&2;

export GIMME_ENV="/home/student/.gimme/envs/go1.10.1.env"
```

Ensure the expected path has been set and exported:

```
student@ckad-1:~/test/sonobuoy$ export PATH=$GOROOT/bin:$GOPATH/bin:$PATH
```

Use the `go` command to download the `sonobuoy` code:

```
student@ckad-1:~/test/sonobuoy$ go get -u -v github.com/heptio/sonobuoy
github.com/heptio/sonobuoy (download)
created GOPATH=/home/student/go; see 'go help gopath'
github.com/heptio/sonobuoy/pkg/buildinfo
<output_omitted>
```

Execute the newly downloaded tool with the `run` option. Review the output. Take note of interesting tests in order to search for particular output in the logs.

```
student@ckad-1:~/test/sonobuoy$ /home/student/go/bin/sonobuoy run
Running plugins: e2e, systemd-logs
INFO[0000] created object
name=heptio-sonobuoy namespace= resource=namespaces
INFO[0000] created object
name=sonobuoy-serviceaccount namespace=heptio-sonobuoy
resource=serviceaccounts
INFO[0000] created object
name=sonobuoy-serviceaccount-heptio-sonobuoy namespace=
resource=clusterrolebindings
INFO[0000] created object
name=sonobuoy-serviceaccount namespace= resource=clusterroles
INFO[0000] created object
name=sonobuoy-config-cm namespace=heptio-sonobuoy resource=configmaps
INFO[0000] created object
name=sonobuoy-plugins-cm namespace=heptio-sonobuoy resource=configmaps
INFO[0000] created object                               name=sonobuoy
namespace=heptio-sonobuoy resource=pods
```

```
INFO[0000] created object
name=sonobuoy-master namespace=heptio-sonobuoy resource=services
```

Check the status of `sonobuoy`. It can take up to an hour to finish on large clusters. On our two-node cluster, it will take about two minutes.

```
student@ckad-1:~/test/sonobuoy$ /home/student/go/bin/sonobuoy status
PLUGIN          STATUS          COUNT
e2e              running         1
systemd_logs     complete       1
systemd_logs     running         1
```

Sonobuoy is still running. Runs can take up to 60 minutes.

Look at the logs. If the tests are ongoing, you will see incomplete logs.

```
student@ckad-1:~/test/sonobuoy$ /home/student/go/bin/sonobuoy logs
namespace="heptio-sonobuoy"
pod="sonobuoy-systemd-logs-daemon-set-e322ef32b0804cd2-d48np"
container="sonobuoy-worker"
time="2018-04-15T20:50:48Z" level=info msg="Waiting for waitfile"
waitfile=/tmp/results/done
time="2018-04-15T20:50:49Z" level=info msg="Detected done file,
transmitting result file" resultFile=/tmp/results/systemd_logs
namespace="heptio-sonobuoy" pod="sonobuoy" container="kube-sonobuoy"
<output_omitted>
```

Change into the `client` directory and look at the tests and results generated:

```
student@ckad-1:~/test/sonobuoy$ cd /home/student/test/sonobuoy/pkg/client/

student@ckad-1:~/test/sonobuoy/pkg/client$ ls
defaults.go  doc.go  example_interfaces_test.go  gen_test.go  logs.go
mode.go      results  run.go
delete.go    e2e.go  gen.go                      interfaces.go
logs_test.go  preflight.go  retrieve.go  status.go

student@ckad-1:~/test/sonobuoy/pkg/client$ cd results/

student@ckad-1:~/test/sonobuoy/pkg/client/results$ ls
doc.go  e2e  junit_utils.go  reader.go  reader_test.go  testdata  types.go

student@ckad-1:~/test/sonobuoy/pkg/client/results$ cd testdata/
```

```
student@ckad-1:~/test/sonobuoy/pkg/client/results/testdata$ ls -l
total 644
-rw-rw-r-- 1 student student 407010 Apr 15 20:43 results-0.10.tar.gz
-rw-rw-r-- 1 student student 32588 Apr 15 20:43 results-0.8.tar.gz
-rw-rw-r-- 1 student student 215876 Apr 15 20:43 results-0.9.tar.gz
```

```
student@ckad-1:~/test/sonobuoy/pkg/client/results/testdata$ tar -xf
results-0.8.tar.gz
```

```
student@ckad-1:~/test/sonobuoy/pkg/client/results/testdata$ ls
config.json  hosts  plugins  resources  results-0.10.tar.gz
results-0.8.tar.gz  results-0.9.tar.gz  serverversion
```

```
student@ckad-1:~/test/sonobuoy/pkg/client/results/testdata$ less \
plugins/e2e/results/e2e.log
Dec 13 17:06:53.480: INFO: Overriding default scale value of zero to 1
Dec 13 17:06:53.481: INFO: Overriding default milliseconds value of zero to
5000
Running Suite: Kubernetes e2e suite
=====
Random Seed: 1513184813 - Will randomize all specs
Will run 1 of 698 specs
```

```
Dec 13 17:06:54.705: INFO: >>> kubeConfig:
Dec 13 17:06:54.707: INFO: Waiting up to 4h0m0s for all (but 0) nodes to be
schedulable
Dec 13 17:06:54.735: INFO: Waiting up to 10m0s for all pods (need at least
0) in namespace 'kube-system' to be running a
nd ready
Dec 13 17:06:54.895: INFO: 14 / 14 pods in namespace 'kube-system' are
running and ready (0 seconds elapsed)
Dec 13 17:06:54.895: INFO: expected 3 pod replicas in namespace
'kube-system', 3 are Running and Ready.
<output_omitted>
```

Find other files which have been generated, and their size:

```
student@ckad-1:~/t.../testdata$ find podlogs/ -exec ls -l {} \;
total 4
drwxr-xr-x 7 student student 4096 Dec 13 17:07 heptio-sonobuoy
total 20
drwxr-xr-x 3 student student 4096 Apr 15 21:28 sonobuoy
```

```

drwxr-xr-x 3 student student 4096 Apr 15 21:28
sonobuoy-e2e-job-b803642b9d884c42
drwxr-xr-x 3 student student 4096 Apr 15 21:28 systemd-logs-btk5c
drwxr-xr-x 3 student student 4096 Dec 13 17:07 systemd-logs-mv5rw
drwxr-xr-x 3 student student 4096 Apr 15 21:28 systemd-logs-rzj5
total 4
drwxr-xr-x 2 student student 4096 Dec 13 17:07 logs
total 8
-rw-r--r-- 1 student student 4568 Dec 13 17:07 kube-sonobuoy.txt
-rw-r--r-- 1 student student 4568 Dec 13 17:07
podlogs/heptio-sonobuoy/sonobuoy/logs/kube-sonobuoy.txt
<output_omitted>

```

Read through some of the tests and output.

```

student@ckad-1:~/test/sonobuoy/pkg/client/results/testdata$ less \
podlogs/heptio-sonobuoy/sonobuoy/logs/kube-sonobuoy.txt

```

Continue to look through tests and results as time permits. There is also an online, graphical scanner. In testing inside GCE, the results were blocked and never returned. You may have different outcome in other environments.